

# **CS 4530 Software Engineering**

## **Module 9: Distributed Systems Principles and Requirements**

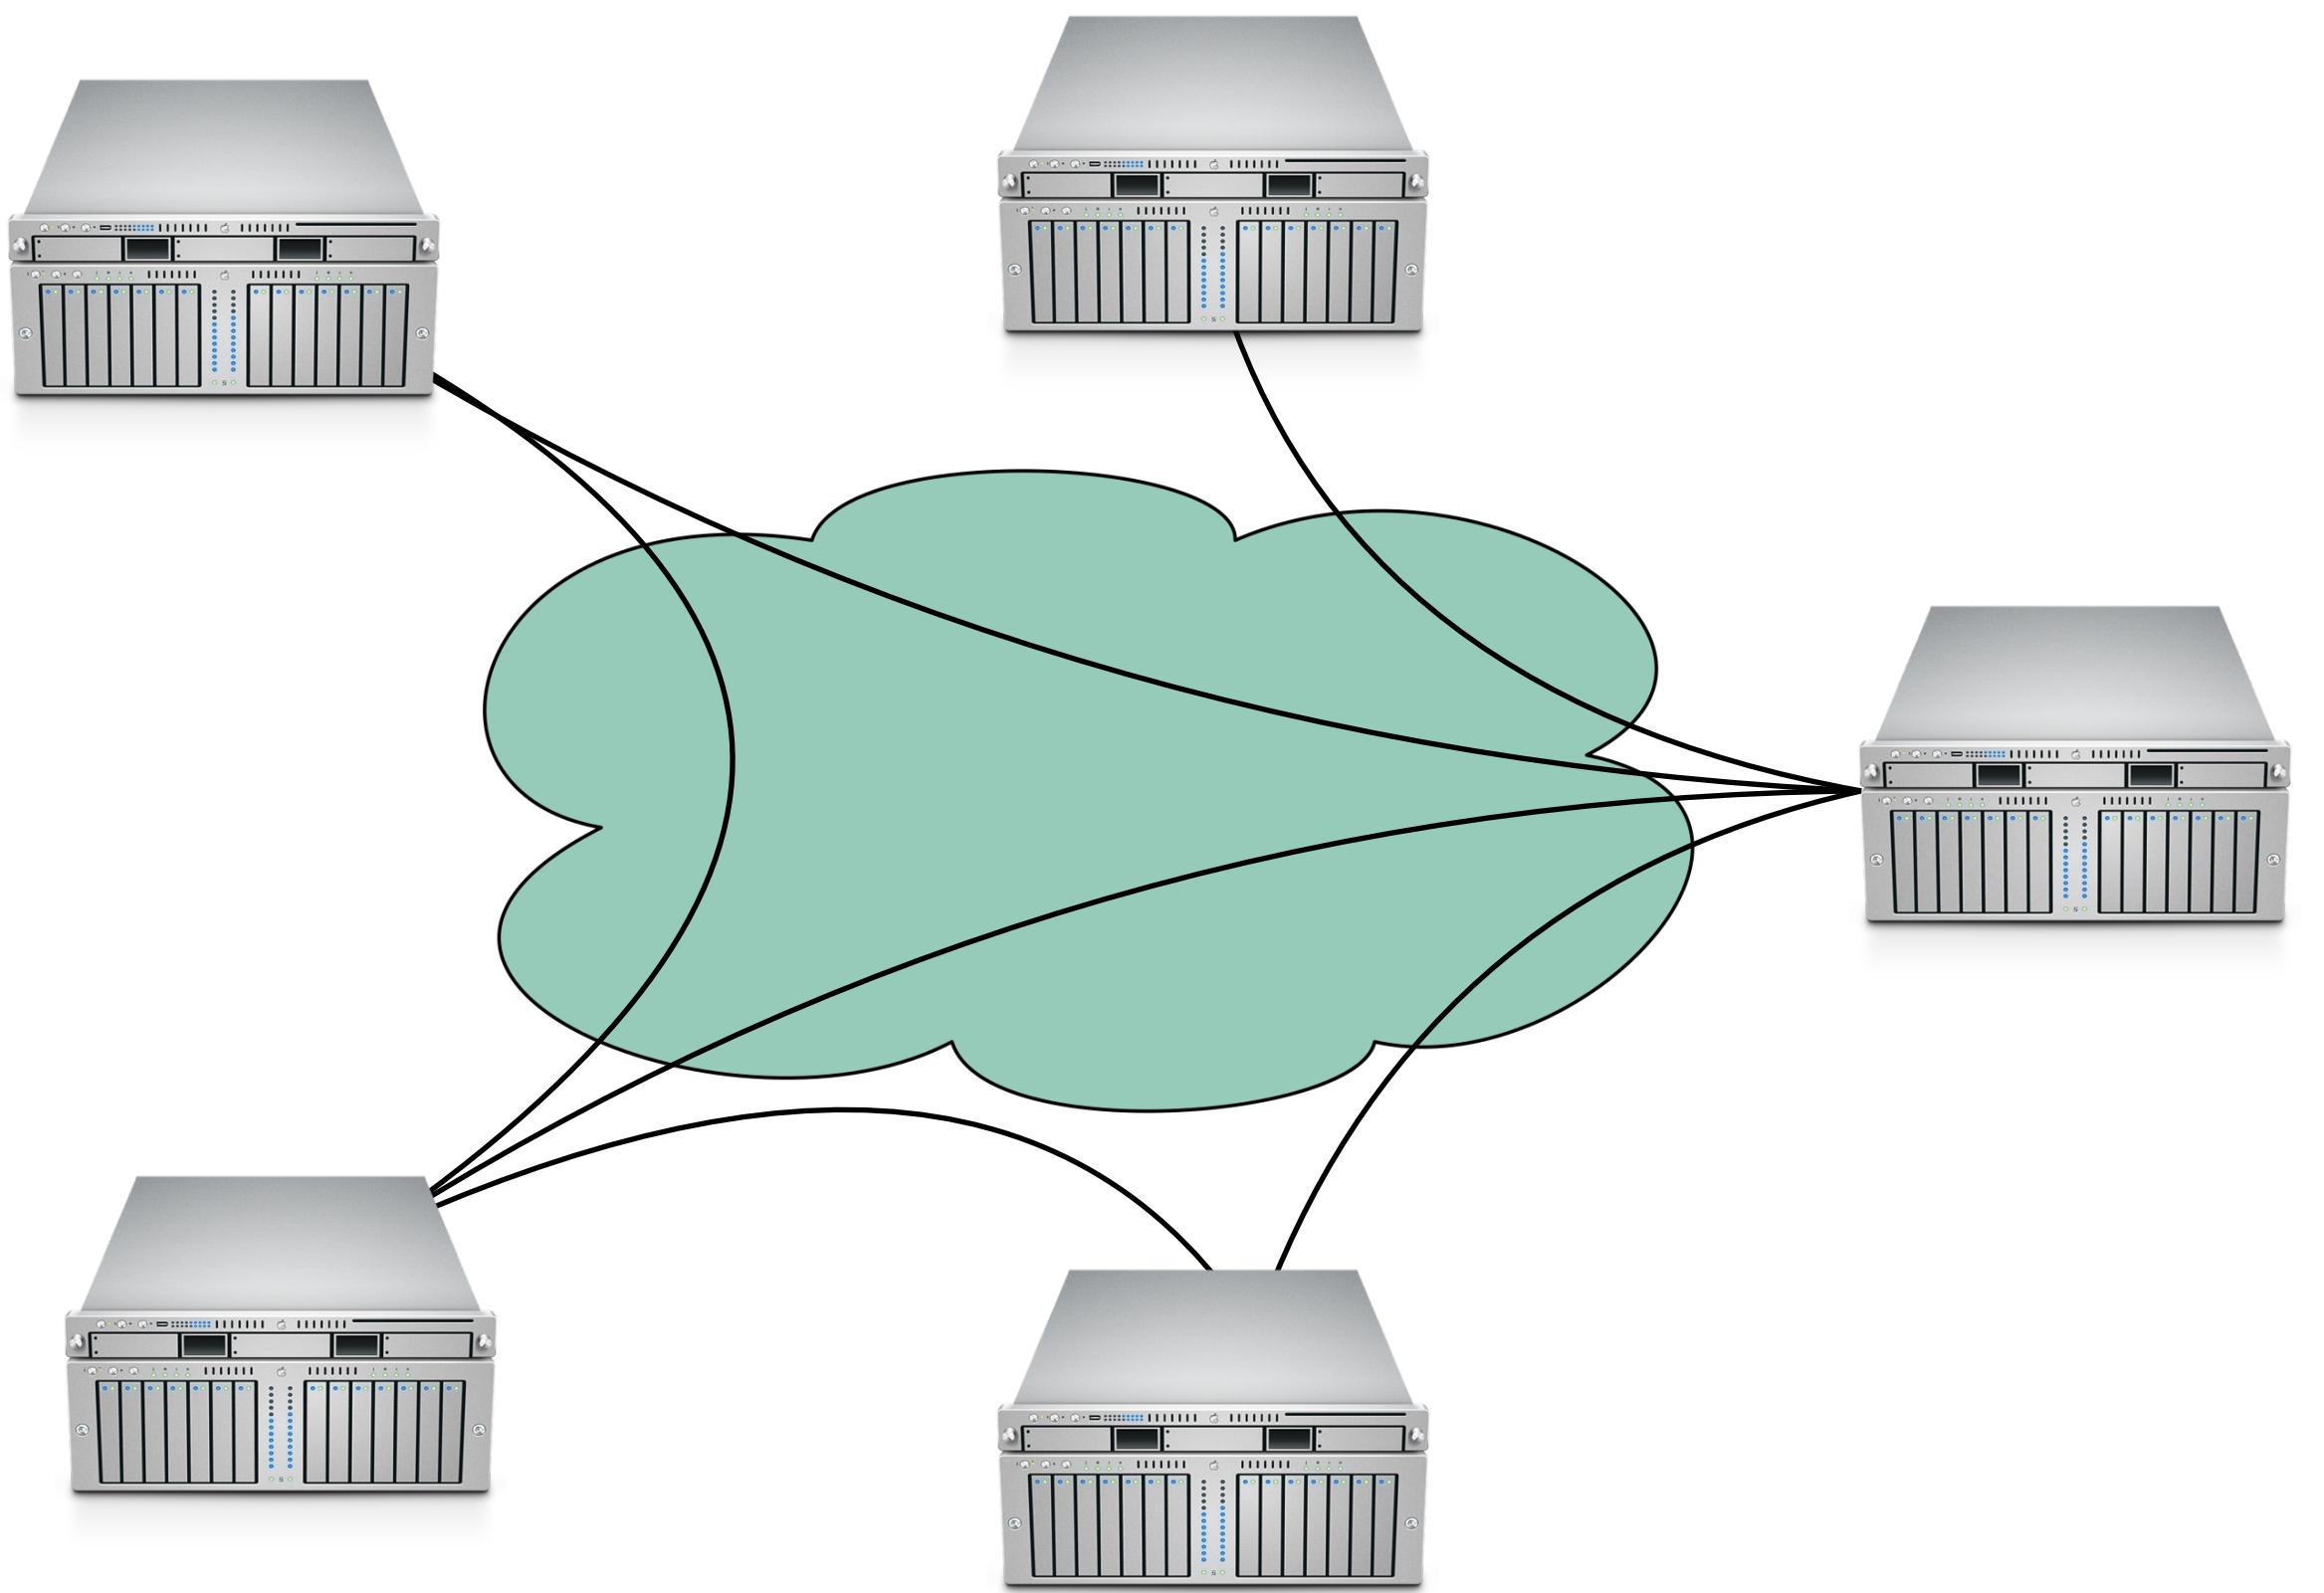
**Jonathan Bell, Adeel Bhutta, Mitch Wand**  
**Khoury College of Computer Sciences**  
**© 2022, released under [CC BY-SA](#)**

# Learning Objectives for this Lesson

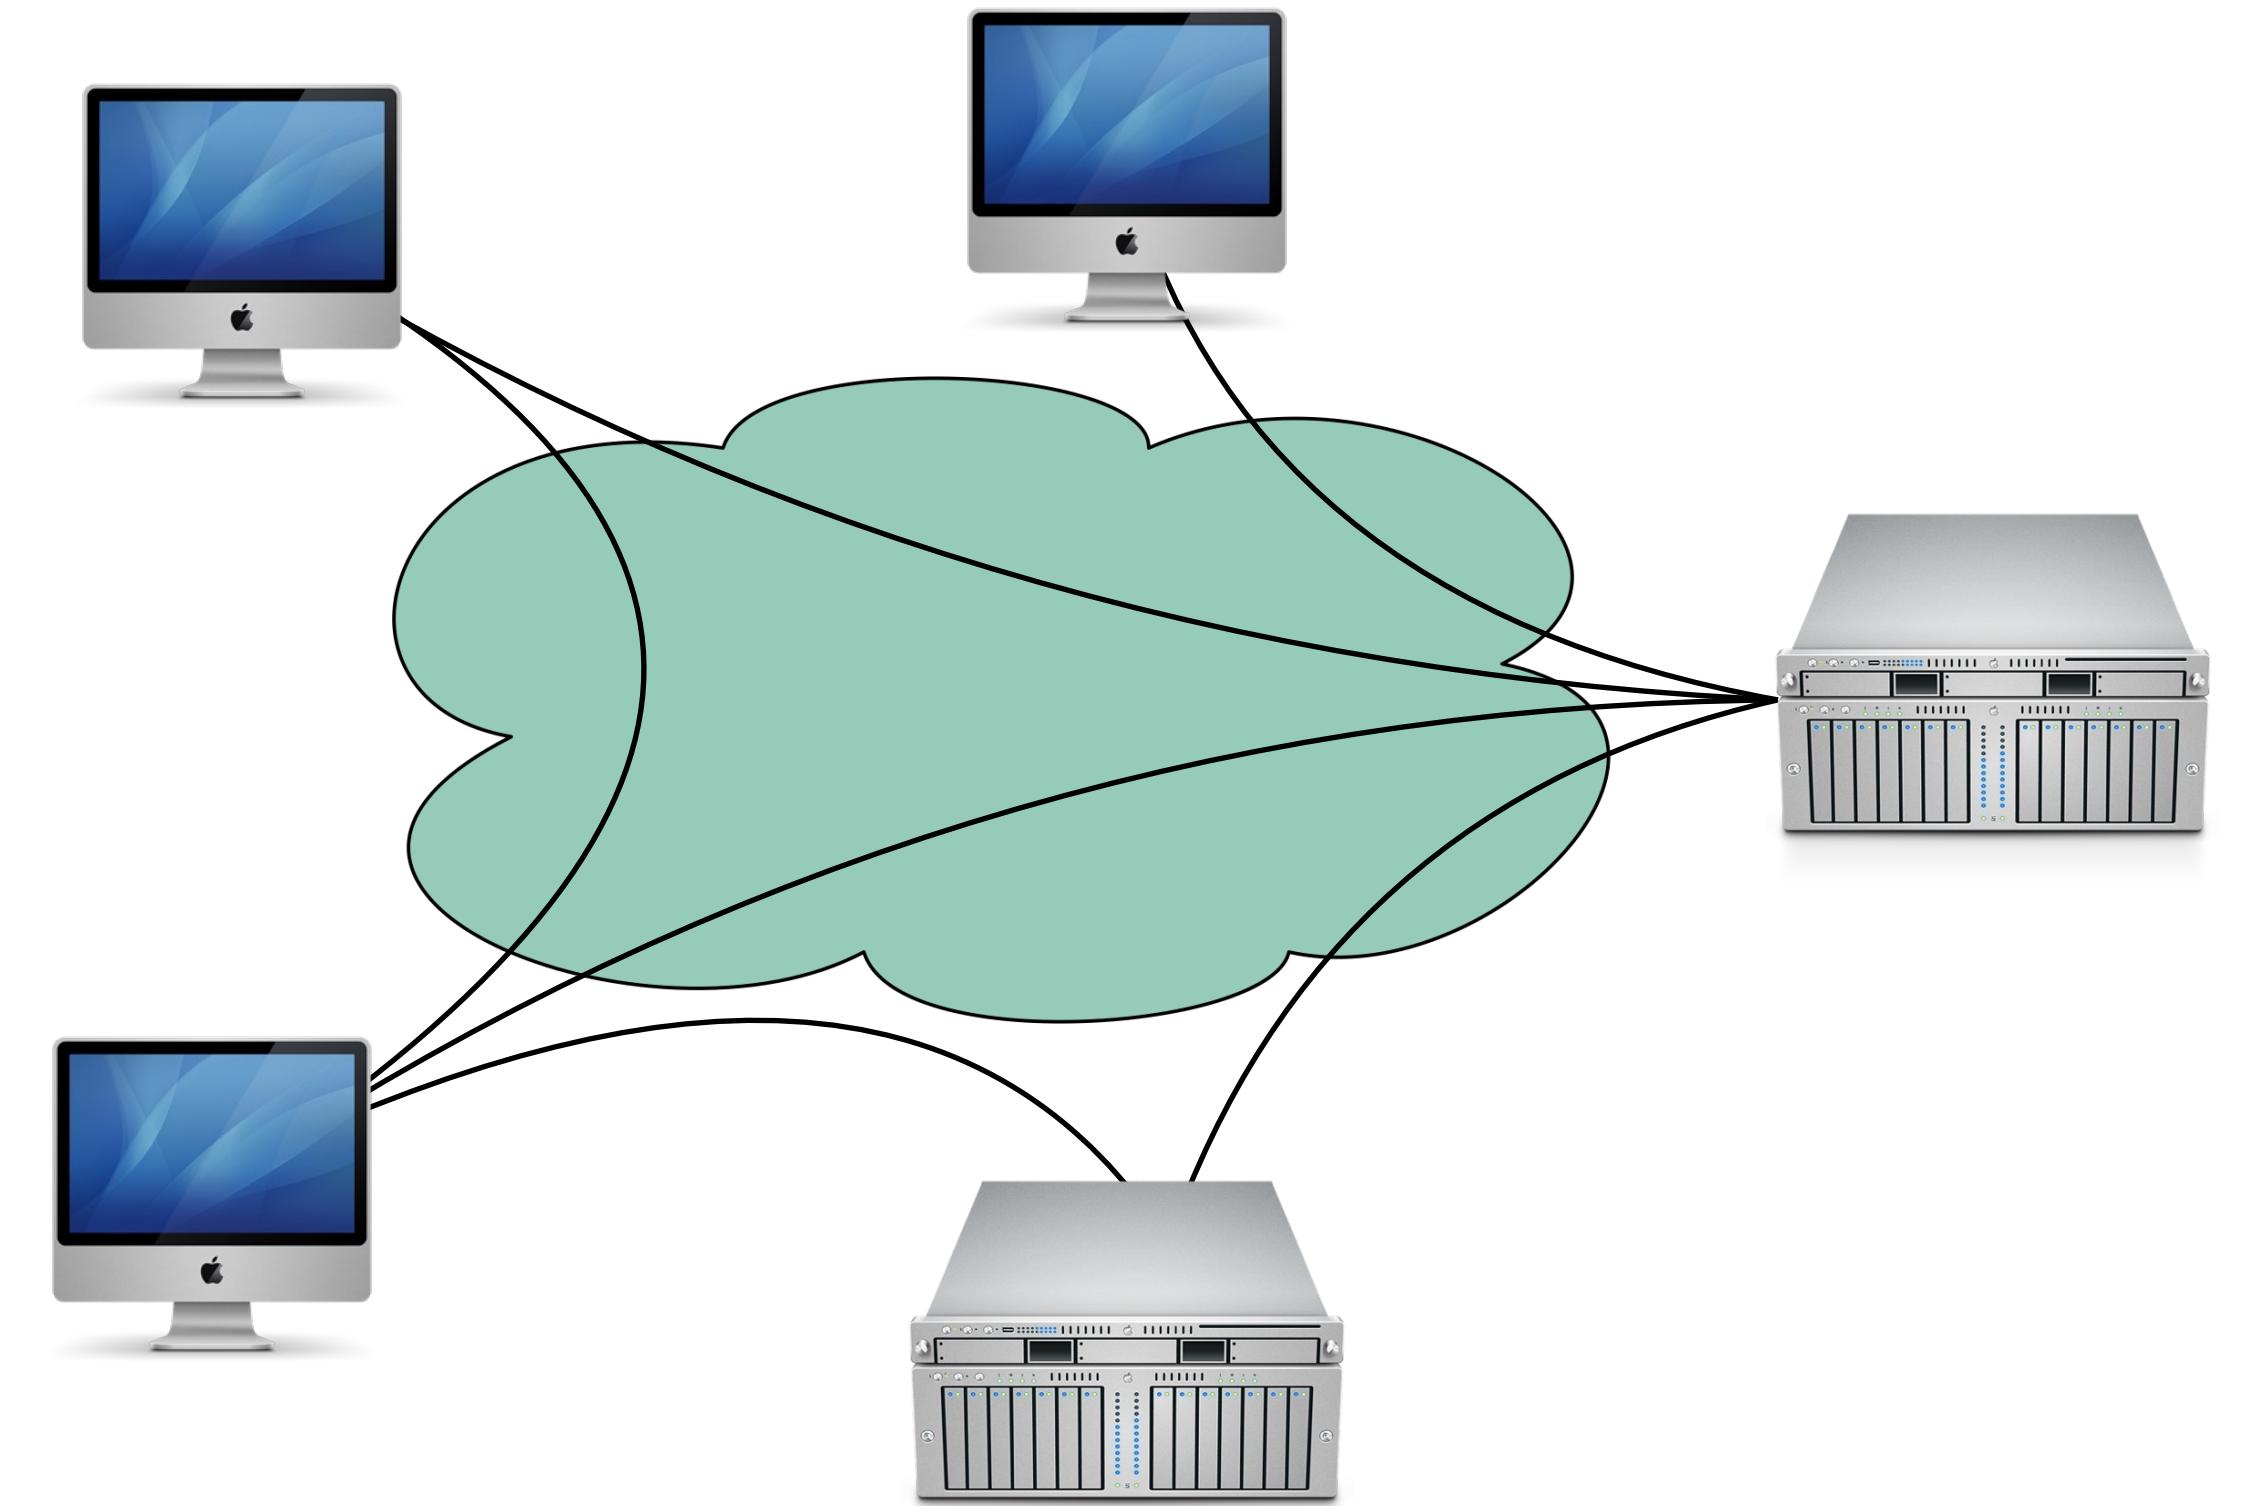
**By the end of this lesson, you should be able to...**

- Describe 5 key goals of distributed systems
- Understand the fundamental constraints of distributed systems
- Understand the roles of replication and partitioning in the context of the DNS infrastructure

# What is a distributed system?



Model:  
Many servers talking through a network



Model:  
Many servers and clients talking through a network

# Why expand to distributed systems?

(Here is why) - performance, availability, fault tolerance

- Scalability
- Performance
- Latency
- Availability
- Fault Tolerance

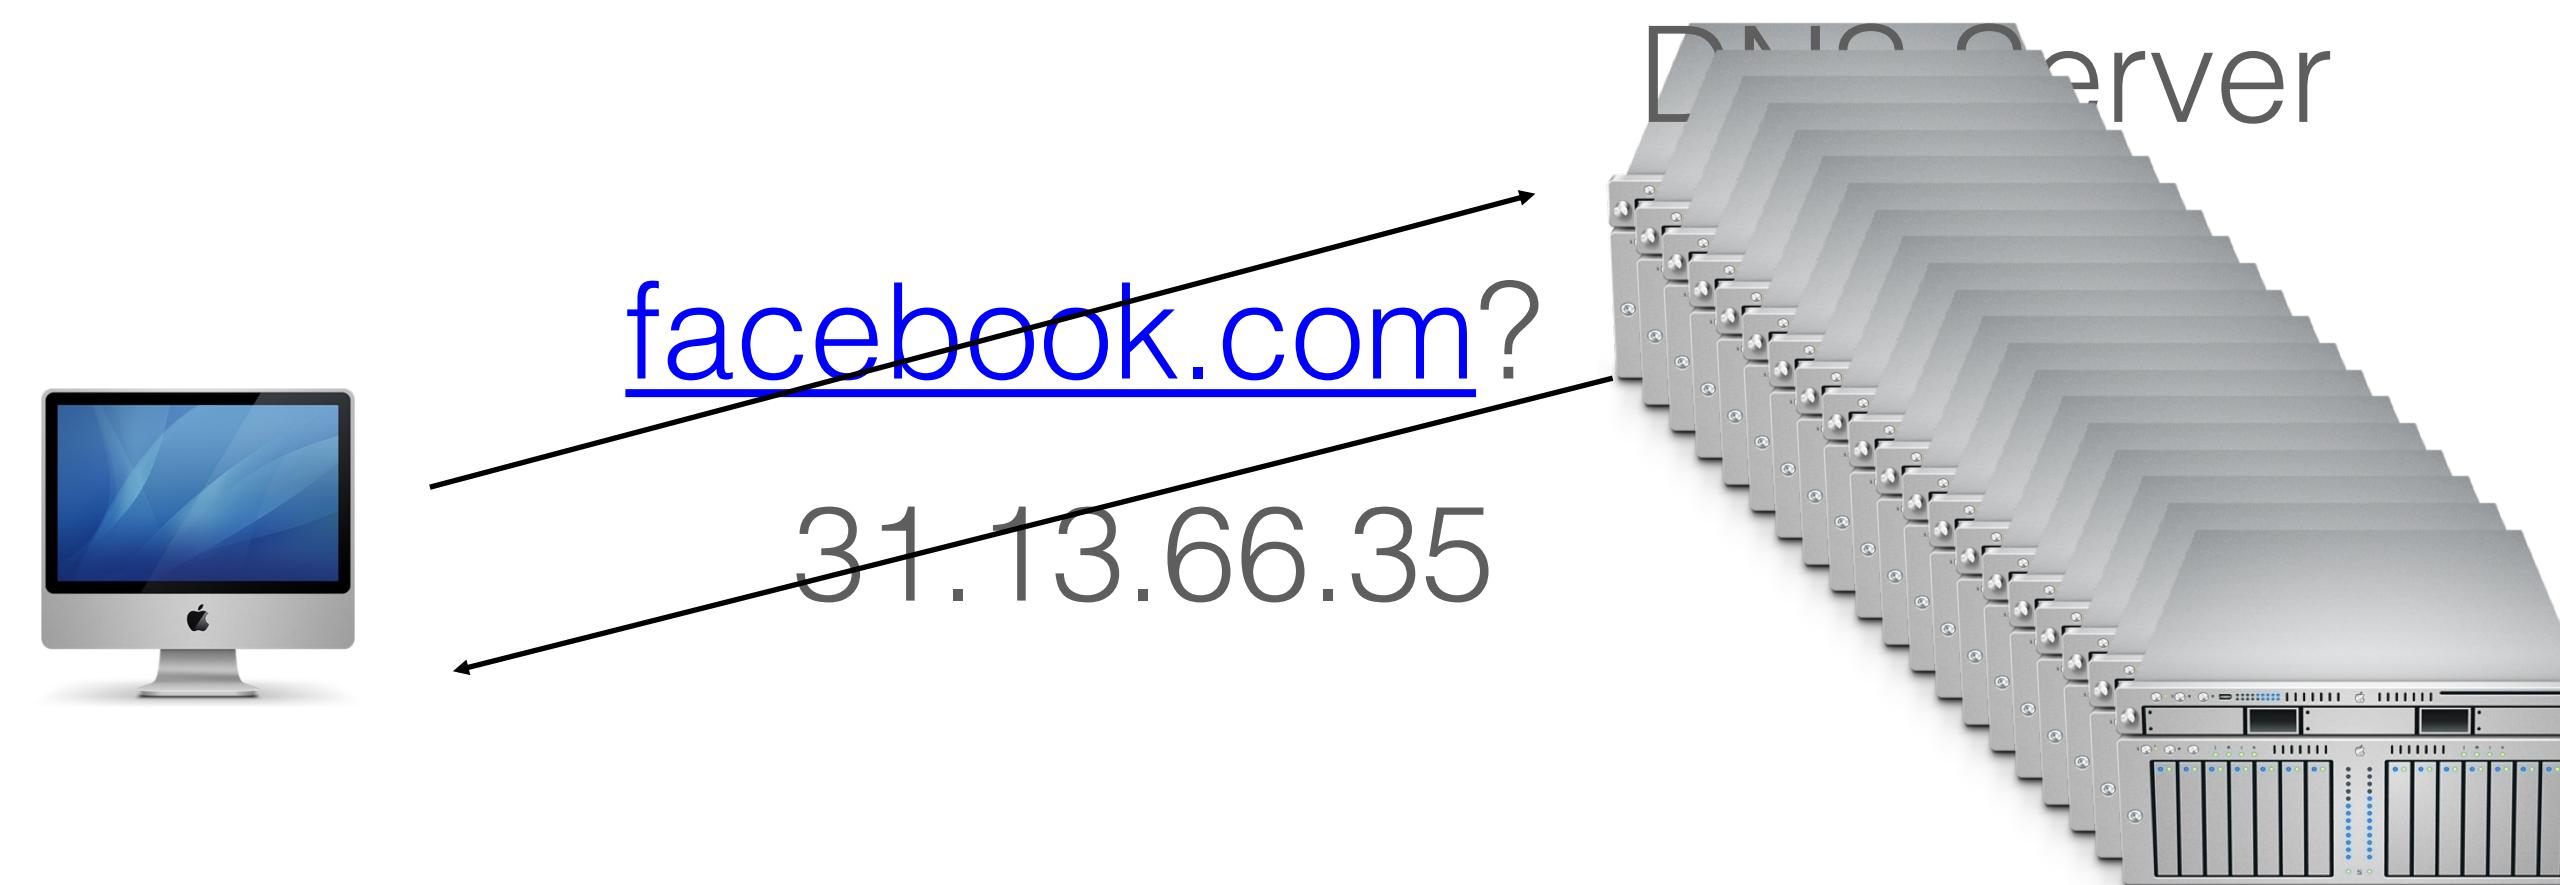
# Example: Domain Name System (DNS)

## Problem Statement

- Nodes (hosts) on a network are identified by IP addresses
- E.g.: 142.251.41.4
- We humans prefer something easier to remember: calendar.google.com, facebook.com, www.khoury.northeastern.edu
- We need to keep a directory of domain names and their addresses
- We also need to make sure everybody gets directed to the correct host

# Example: Domain Name System (DNS)

- Need to handle millions of DNS queries per second
- Not immediately obvious how to scale: how do we maintain replication, some measure of consistency?



# Domain Name System

- Strawman solution: Use a Local file
  - Keep local copy of mapping from all hosts to all IPs (e.g., /etc/hosts)
  - Hosts change IPs regularly: Download file frequently
  - Lot of constant internet bandwidth use
  - IPv4 space is now full
    - 32-bits: 4,294,967,296 addresses
    - At 1 byte per address, file would be 4GB
    - Not a lot of disk space (now, DNS introduced in the late 80s)

# Domain Name System

- Strawman solution: Use a Local file
  - Keep local copy of mapping from all hosts to all IPs (e.g., /etc/hosts)
  - Hosts change IPs regularly: Download file frequently
  - Lot of constant internet bandwidth use
  - IPv4 space is now full
    - 32-bits: 4,294,967,296 addresses
    - At 1 byte per address, file would be 4GB
  - Not a lot of c



We need 200x of these  
to hold 4GB: \$270K+

**Inflation Calculator**

If in  (enter year)

I purchased an item for \$

then in  (enter year)

that same item would cost: **\$1,391.65**

Cumulative rate of inflation: **98.8%**

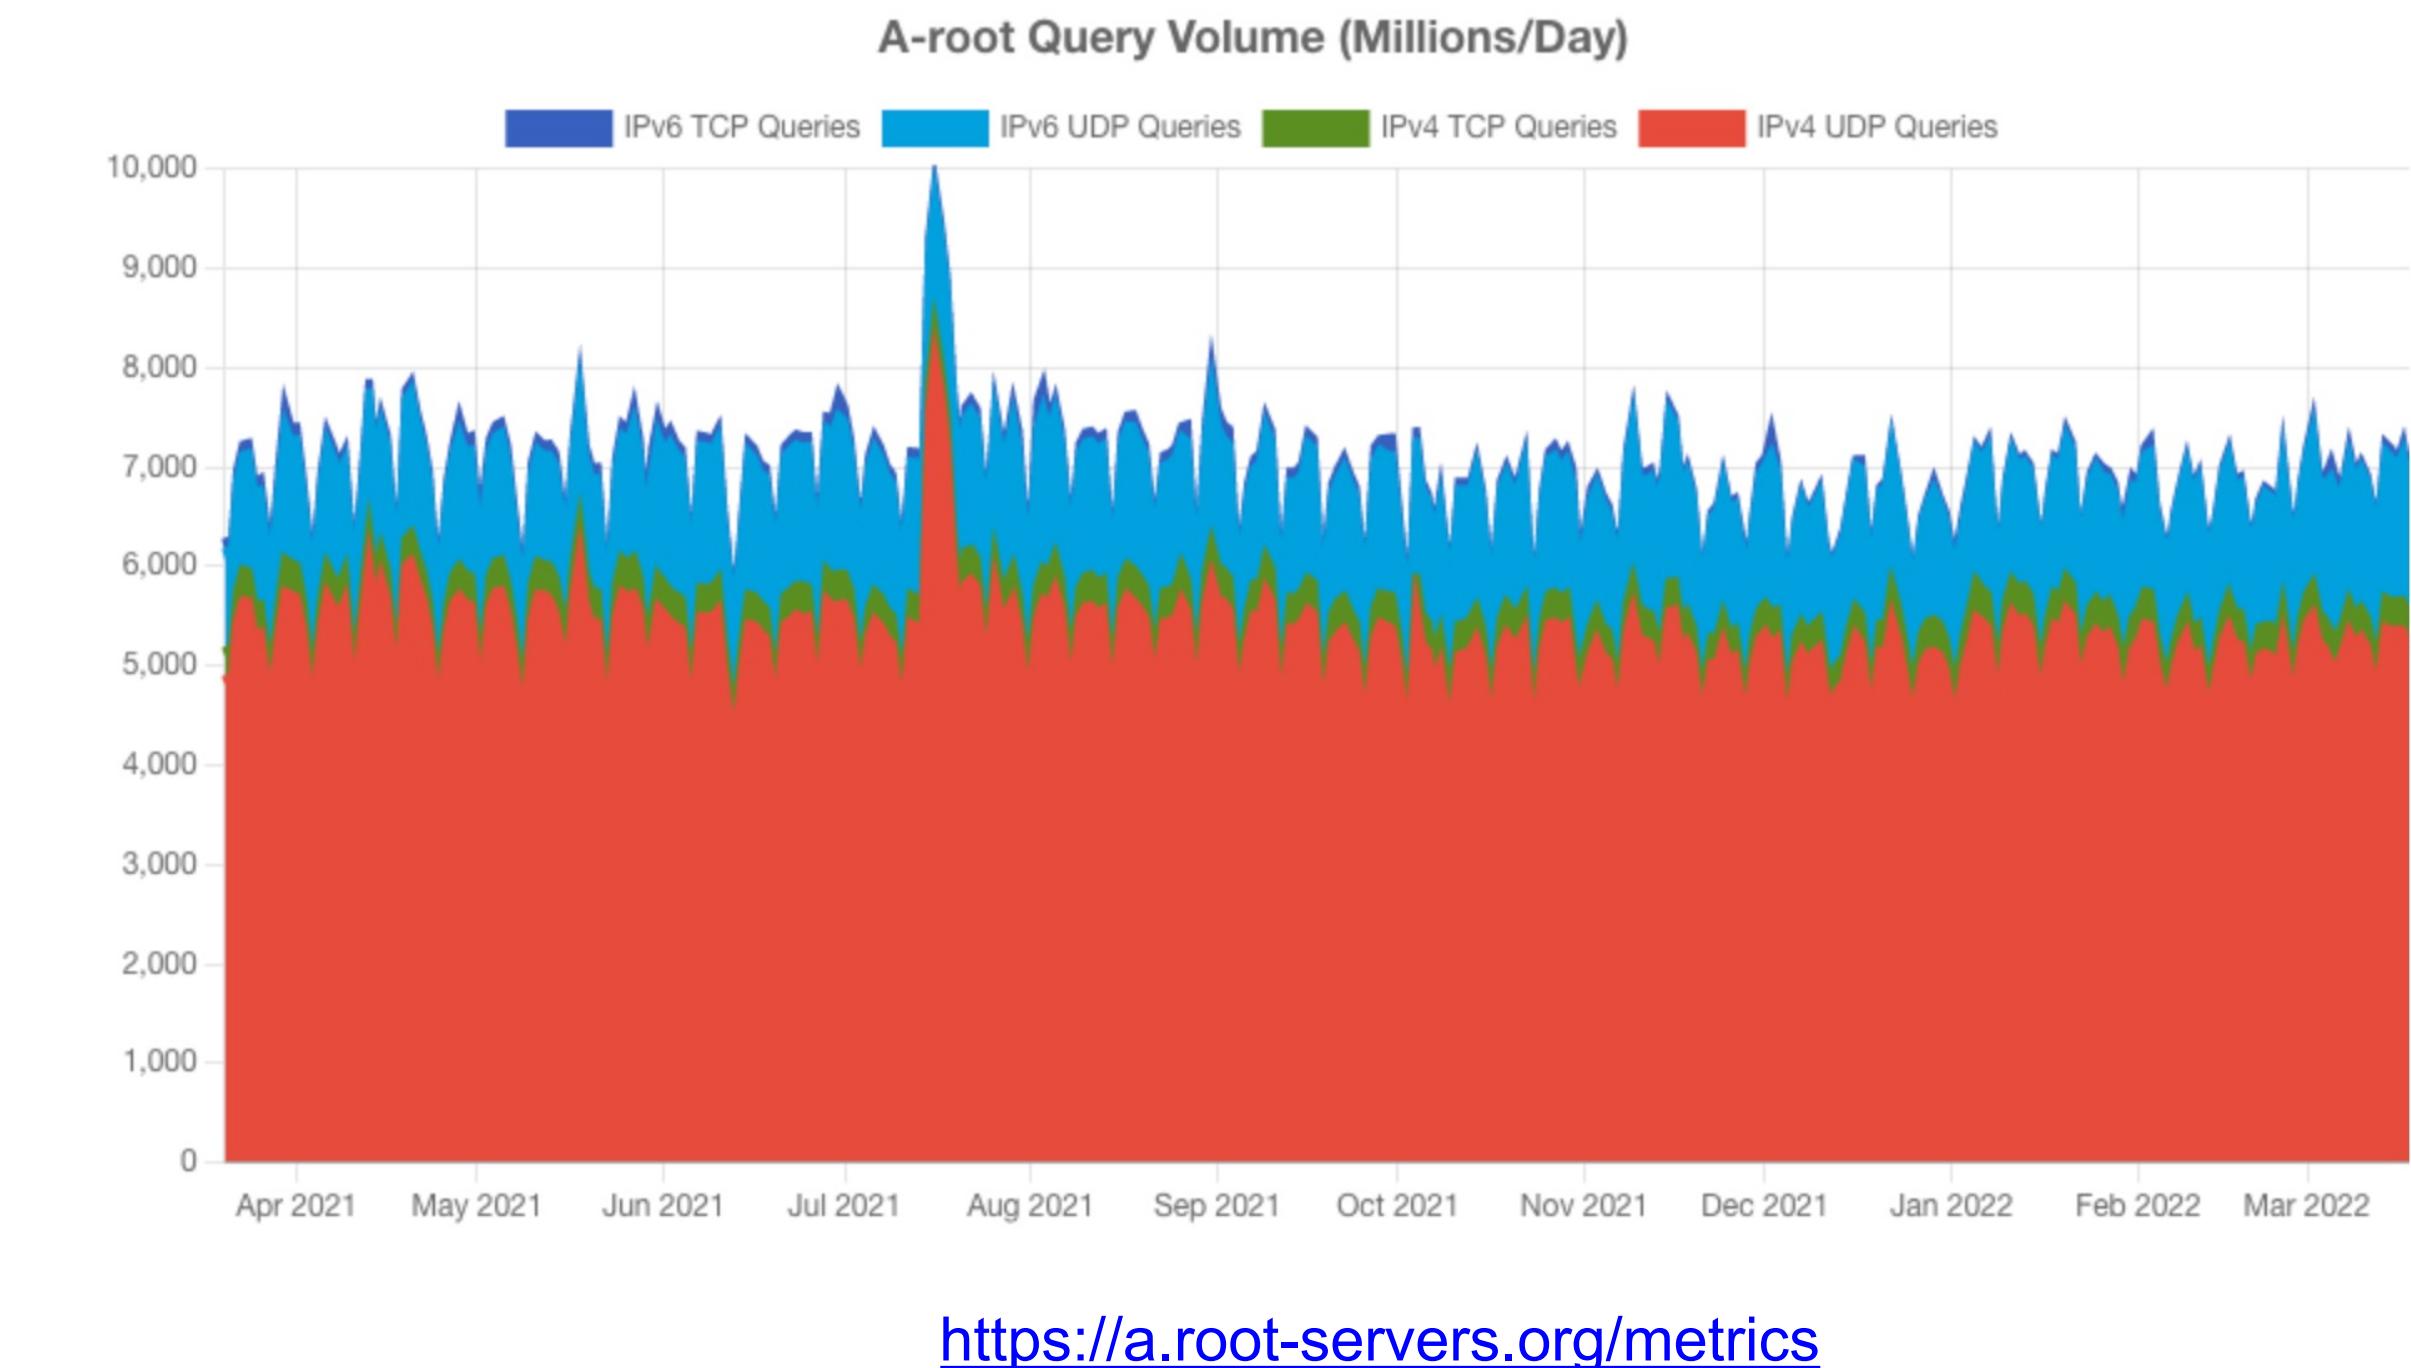
**CALCULATE**

# Domain Name System

- Strawman solution: Use a Local file
  - Keep local copy of mapping from all hosts to all IPs (e.g., /etc/hosts)
  - Hosts change IPs regularly: Download file frequently
  - Lot of constant internet bandwidth use
  - IPv4 space is now full
    - 32-bits: 4,294,967,296 addresses
    - At 1 byte per address, file would be 4GB
    - Not a lot of disk space (now, DNS introduced in the late 80s)
    - But a lot of constant internet bandwidth
  - More names than IPs
    - Aliases
  - **Not scalable!**

# Domain Name System

- Another Strawman: Well-known centralized server. All requests made to this server:
  - Single point of failure
  - Bottleneck for throughput and access time
  - Bottleneck for administration (adding/changing records?)
  - **Ultimately, not scalable!**



# DNS as a distributed system

- We need a **scalable** solution
  - New hosts keep being added
  - Number of users increases
  - Need to maintain speed/responsiveness
- We need our service to be **available** and **fault tolerant**
  - It is a crucial basic service
  - A problematic node shouldn't “crash the internet”
  - Reads are more important than writes: far more queries to resolve records than to update them
- Global in scope
  - Domain names mean the same thing everywhere

# Distributed Systems Goals

- **Scalability**
- Performance
- Latency
- Availability
- Fault Tolerance

“the ability of a system, network, or process, to handle a growing amount of work in a capable manner or its ability to be enlarged to accommodate that growth.”

# Distributed Systems Scale “Horizontally”

- “Vertical” scaling: add more resources to existing server
  - Faster CPUs, more CPU cores, more RAM, more storage
  - Ineffective once: Clock speed plateaus; difficult to write applications that utilize 256 CPU cores (adding 2TB RAM to a server can often help)
- “Horizontal” scaling: add more servers
  - Rely on “commodity” servers rather than state-of-the-art hardware
  - Allows for dynamic addition of resources as needed by load

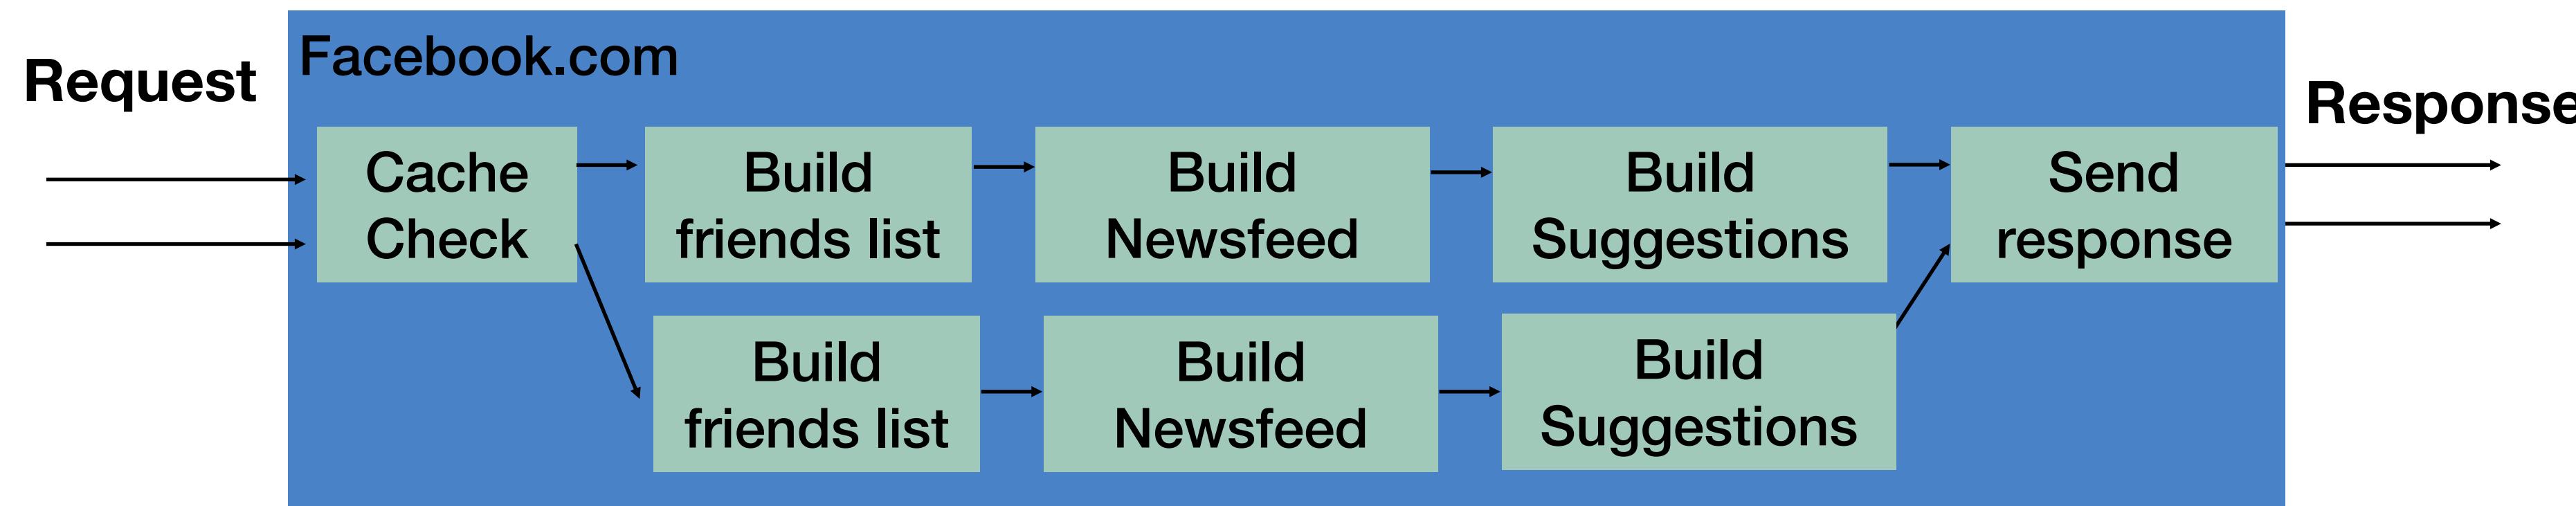
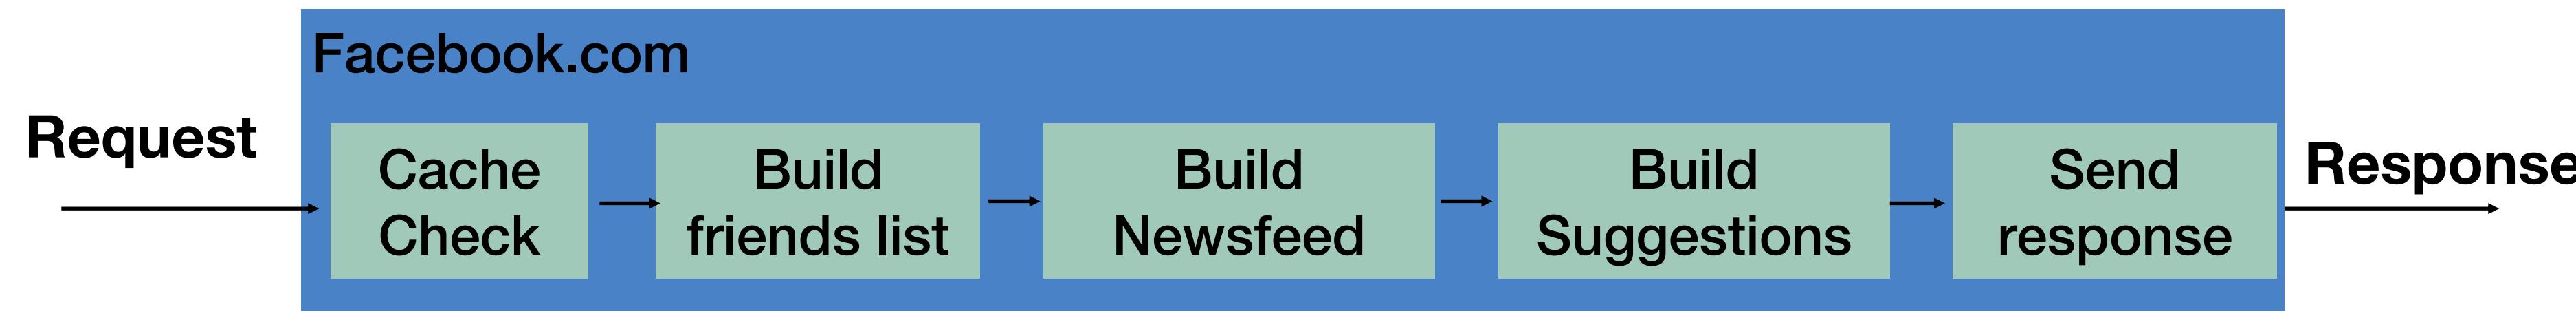
# Distributed Systems Goals

- Scalability
- **Performance**
- Latency
- Availability
- Fault Tolerance

“is characterized by the amount of useful work accomplished by a computer system compared to the time and resources used.”

# Improve Throughput With Concurrency

Throughput: total requests that can be processed per unit-time



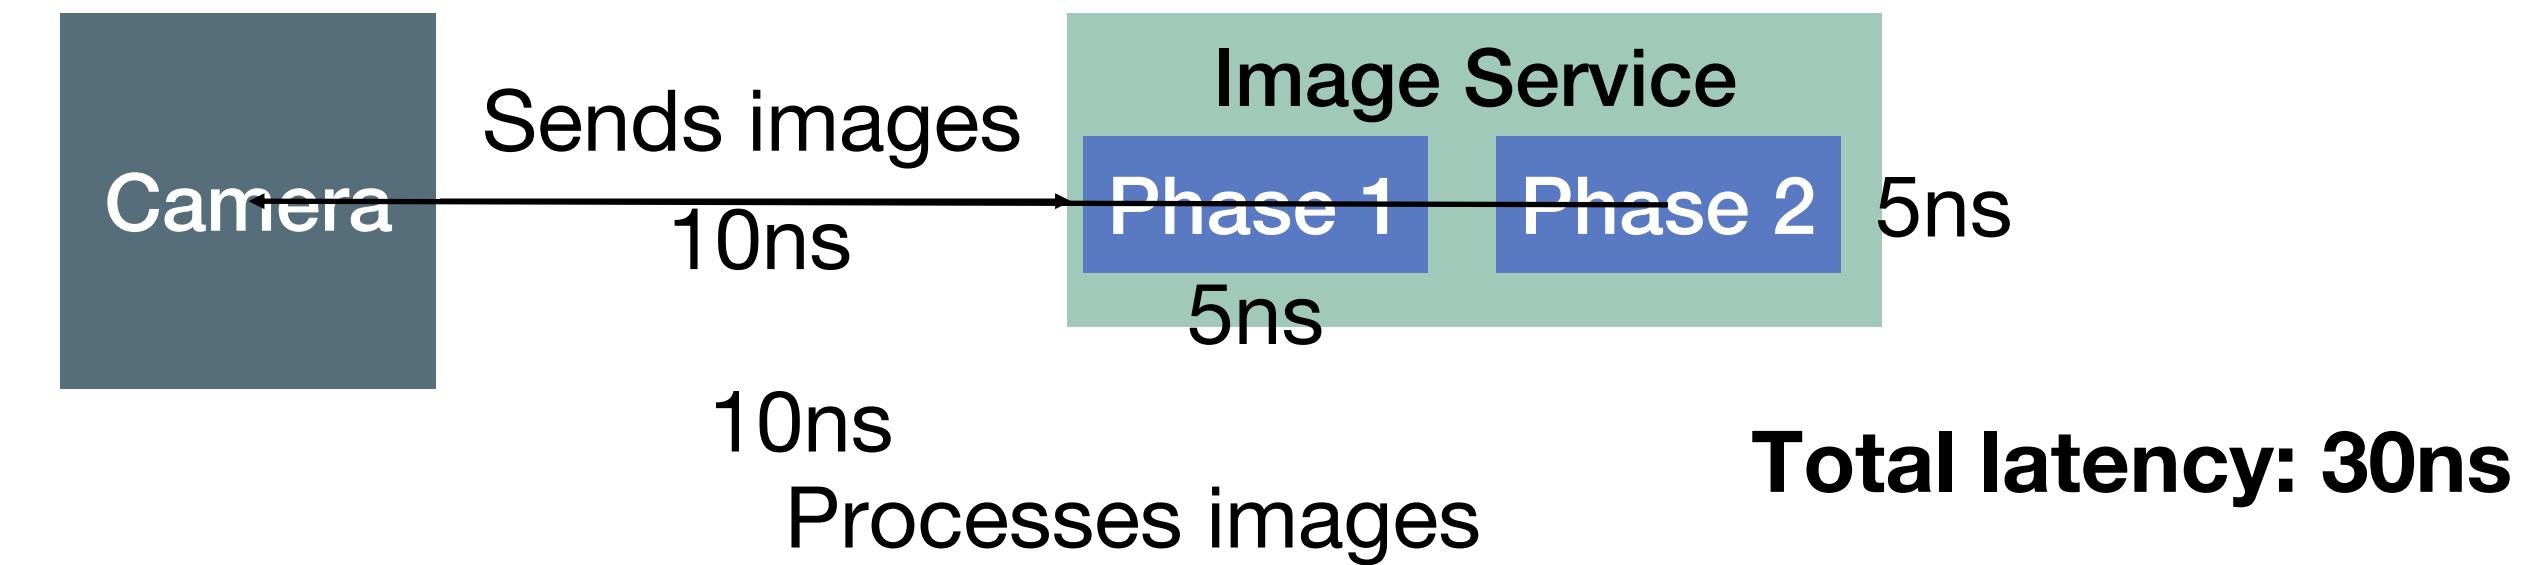
# Distributed Systems Goals

- Scalability
- Performance
- **Latency**
- Availability
- Fault Tolerance

“The state of being latent; delay, a period between the initiation of something and the it becoming visible.”

# Latency: Delay in Receiving a Response

- In client/server model, latency is simply: time between client sending request and receiving response
- What contributes to latency?
- Adding pipelined components -> latency is cumulative



# Improving Latency with Distributed Systems

- Often more challenging than increasing throughput
  - Examples:
    - Physical - Speed of light (network transmissions over long distances)
    - Algorithmic - Looking up an item in a hash table is limited by hash function
    - Economic - Adding more RAM gets expensive
  - Distributed systems can reduce latency by moving the data/server closer to the user

# Networks Introduce Significant Latency

- We still haven't surpassed the speed of light: ~1 CPU cycle  $\approx$  10cm of transmission distance
- Chicago <-> NYC is a valuable link due to commodities markets in Chicago and equities markets in NYC
- High frequency trading created a market for private, lower-latency alternatives to public internet
- Extremely expensive, not a general solution (and still 8.5ms!)



# Distributed Systems Goals

- Scalability
- Performance
- Latency
- **Availability**
- Fault Tolerance

“the proportion of time a system is in a functioning condition. If a user cannot access the system, it is said to be unavailable.”

$$\text{Availability} = \text{uptime} / (\text{uptime} + \text{downtime}).$$

Often measured in “nines”

Availability %	Downtime/year
90%	>1 month
99%	< 4 days
99.9%	< 9 hours
99.99%	<1 hour
99.999%	5 minutes
99.9999%	31 seconds

# Distributed Systems Challenges: Availability

## More machines, more problems

- Say there's a 1% chance of having some hardware failure occur to a machine in a given month (power supply burns out, hard disk crashes, etc)
- Now I have 10 machines
  - Probability(at least one fails during the month) =  $1 - \text{Probability}(\text{no machine fails}) = 1 - (1 - .01)^{10} = 10\%$
- 100 machines -> 63% chance that at least one fails
- 200 machines -> 87% chance that at least one fails (!)
- Implication: System as a whole must tolerate component failures

# Distributed Systems Goals

- Scalability
- Performance
- Latency
- Availability
- **Fault Tolerance**

Disks fail

Power supplies fail

“ability of a system to behave in a well-defined manner once faults occur”

## What kind of faults?

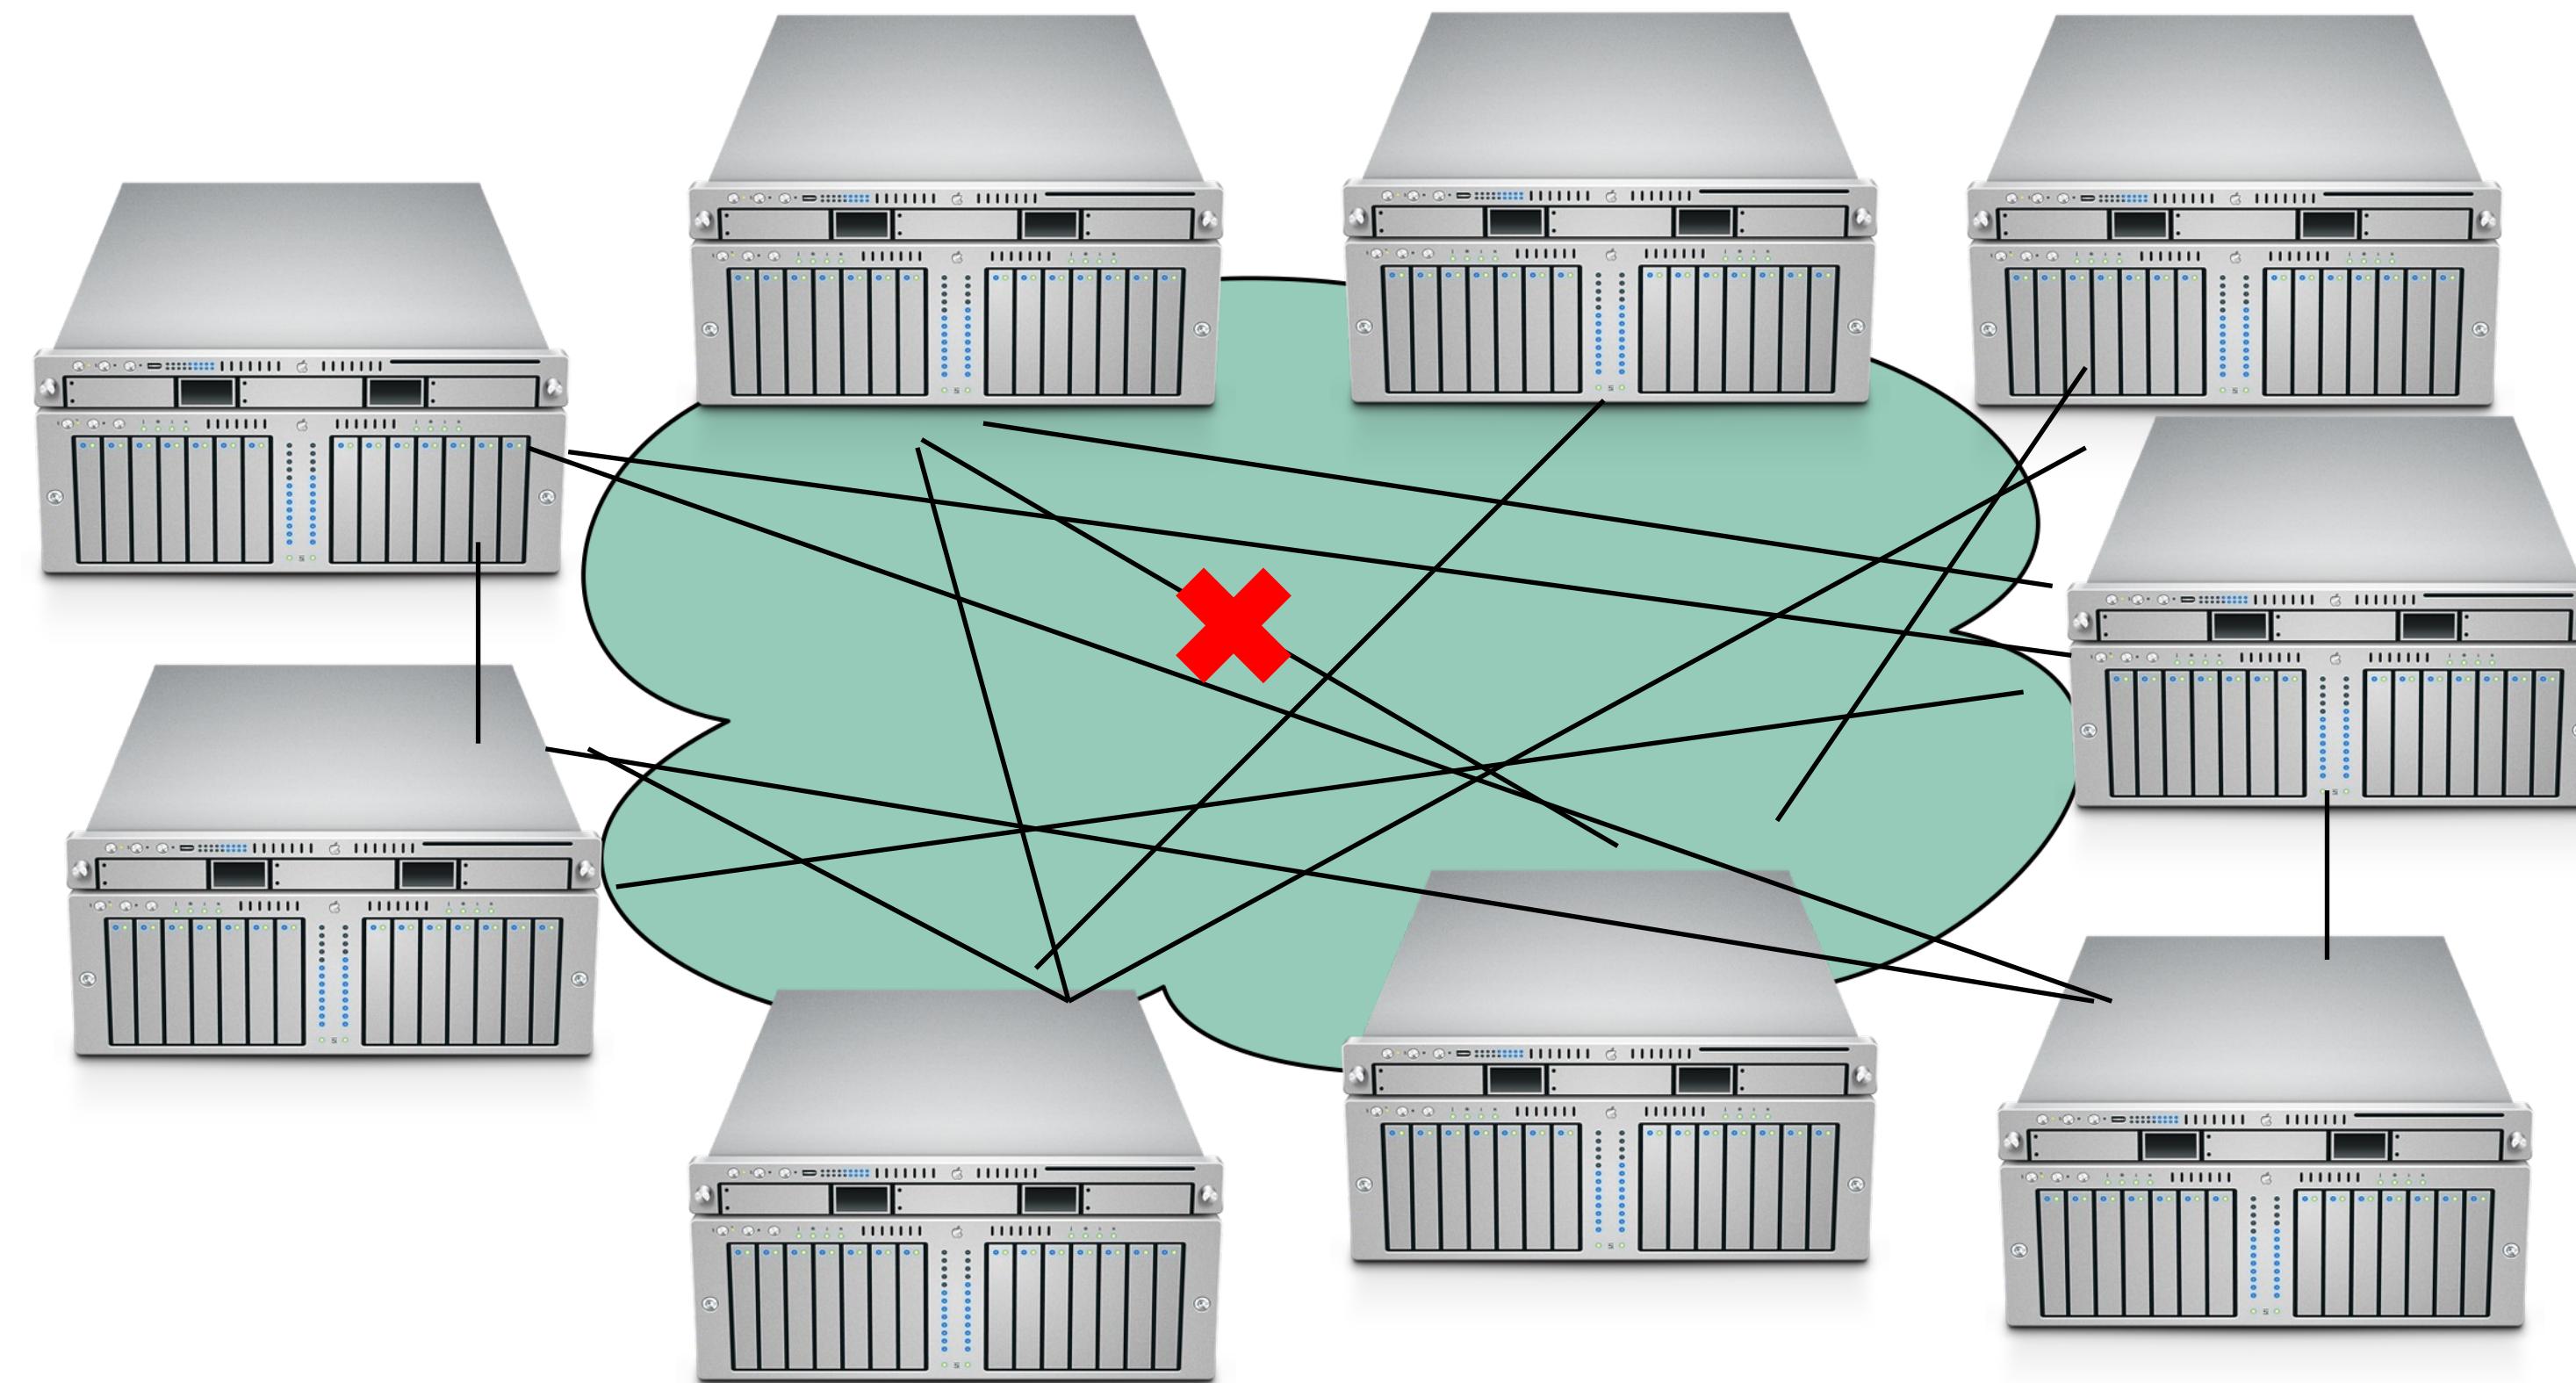
Networking fails

Security breached

Power goes out      Datacenter goes offline

# Distributed Systems Challenges: Performance, Availability

Number of nodes + distance between them



# Challenge: Distributed Systems Rely on Networks

- Can not expect network to be a perfect analog for communication within a single computer because:
  - Speed of light (1 foot/nanosecond)
  - Communication links exist in uncontrolled/hostile environments
  - Communication links may be bandwidth limited (tough to reach even 100MB/sec)
- In contrast to a single computer, where:
  - Distances are measured in mm, not feet
  - Physical concerns can be addressed all at once
  - Bandwidth is plentiful (easily GB/sec)

# And still more challenges

## We still rely on other administrators, who are not infallible

**Amazon Web Services  
outage takes a portion of  
the internet down with it**

Zack Whittaker

@zackwhittaker / 12:32 PM EST • November 25, 2020



 **Image Credits:** David Becker / Getty Images

Amazon Web Services is currently having an outage, taking a chunk of the internet down with it.

Several AWS services were experiencing problems as of early Wednesday, according to [its status page](#). That means any app, site or service that relies on AWS might also be down, too. (As I found out the hard way this morning when



Comment

Summary of the Amazon Kinesis Event in the Northern Virginia (US-EAST-1) Region

**November, 25th 2020**

We wanted to provide you with some additional information about the service disruption that occurred in the Northern Virginia (US-EAST-1) Region on November 25th, 2020.

Amazon Kinesis enables real-time processing of streaming data. In addition to its direct use by customers, Kinesis is used by several other AWS services. These services also saw impact during the event. The trigger, though not root cause, for the event was a relatively small addition of capacity that began to be added to the service at 2:44 AM PST, finishing at 3:47 AM PST. Kinesis has a large number of "back-end" cell-clusters that process streams. These are the workhorses in Kinesis, providing distribution, access, and scalability for stream processing. Streams are spread across the back-end through a sharding mechanism owned by a "front-end" fleet of servers. A back-end cluster owns many shards and provides a consistent scaling unit and fault-isolation. The front-end's job is small but important. It handles authentication, throttling, and request-routing to the correct stream-shards on the back-end clusters.

The capacity addition was being made to the front-end fleet. Each server in the front-end fleet maintains a cache of information, including membership details and shard ownership for the back-end clusters, called a shard-map. This information is obtained through calls to a microservice vending the membership information, retrieval of configuration information from DynamoDB, and continuous processing of messages from other Kinesis front-end servers. For the latter communication, each front-end server creates operating system threads for each of the other servers in the front-end fleet. Upon any addition of capacity, the servers that are already operating members of the fleet will learn of new servers joining and establish the appropriate threads. It takes up to an hour for any existing front-end fleet member to learn of new participants.

At 5:15 AM PST, the first alarms began firing for errors on putting and getting Kinesis records. Teams engaged and began reviewing logs. While the new capacity was a suspect, there were a number of errors that were unrelated to the new capacity and would likely persist even if the capacity were to be removed. Still, as a precaution, we began removing the new capacity while researching the other errors. The diagnosis work was slowed by the variety of errors observed. We were seeing errors in all aspects of the various calls being made by existing and new members of the front-end fleet, exacerbating our ability to separate side-effects from the root cause. At 7:51 AM PST, we had narrowed the root cause to a couple of candidates and determined that any of the most likely sources of the problem would require a full restart of the front-end fleet, which the Kinesis team knew would be a long and careful process. The resources within a front-end server that are used to populate the shard-map compete with the resources that are used to process incoming requests. So, bringing front-end servers back online too quickly would create contention between these two needs and result in very few resources being available to handle incoming requests, leading to increased errors and request latencies. As a result, these slow front-end servers could be deemed unhealthy and removed from the fleet, which in turn, would set back the recovery process. All of the candidate solutions involved changing every front-end server's configuration and restarting it. While the leading candidate (an issue that seemed to be creating memory pressure) looked promising, if we were wrong, we would double the recovery time as we would need to apply a second fix and restart again. To speed restart, in parallel with our investigation, we began adding a configuration to the front-end servers to obtain data directly from the authoritative metadata store rather than from front-end server neighbors during the bootstrap process.

At 9:39 AM PST, we were able to confirm a root cause, and it turned out this wasn't driven by memory pressure. Rather, the new capacity had caused all of the servers in the fleet to exceed the maximum number of threads allowed by an operating system configuration. As this limit was being exceeded,

# System Design Follows Requirements

## Domain Name System

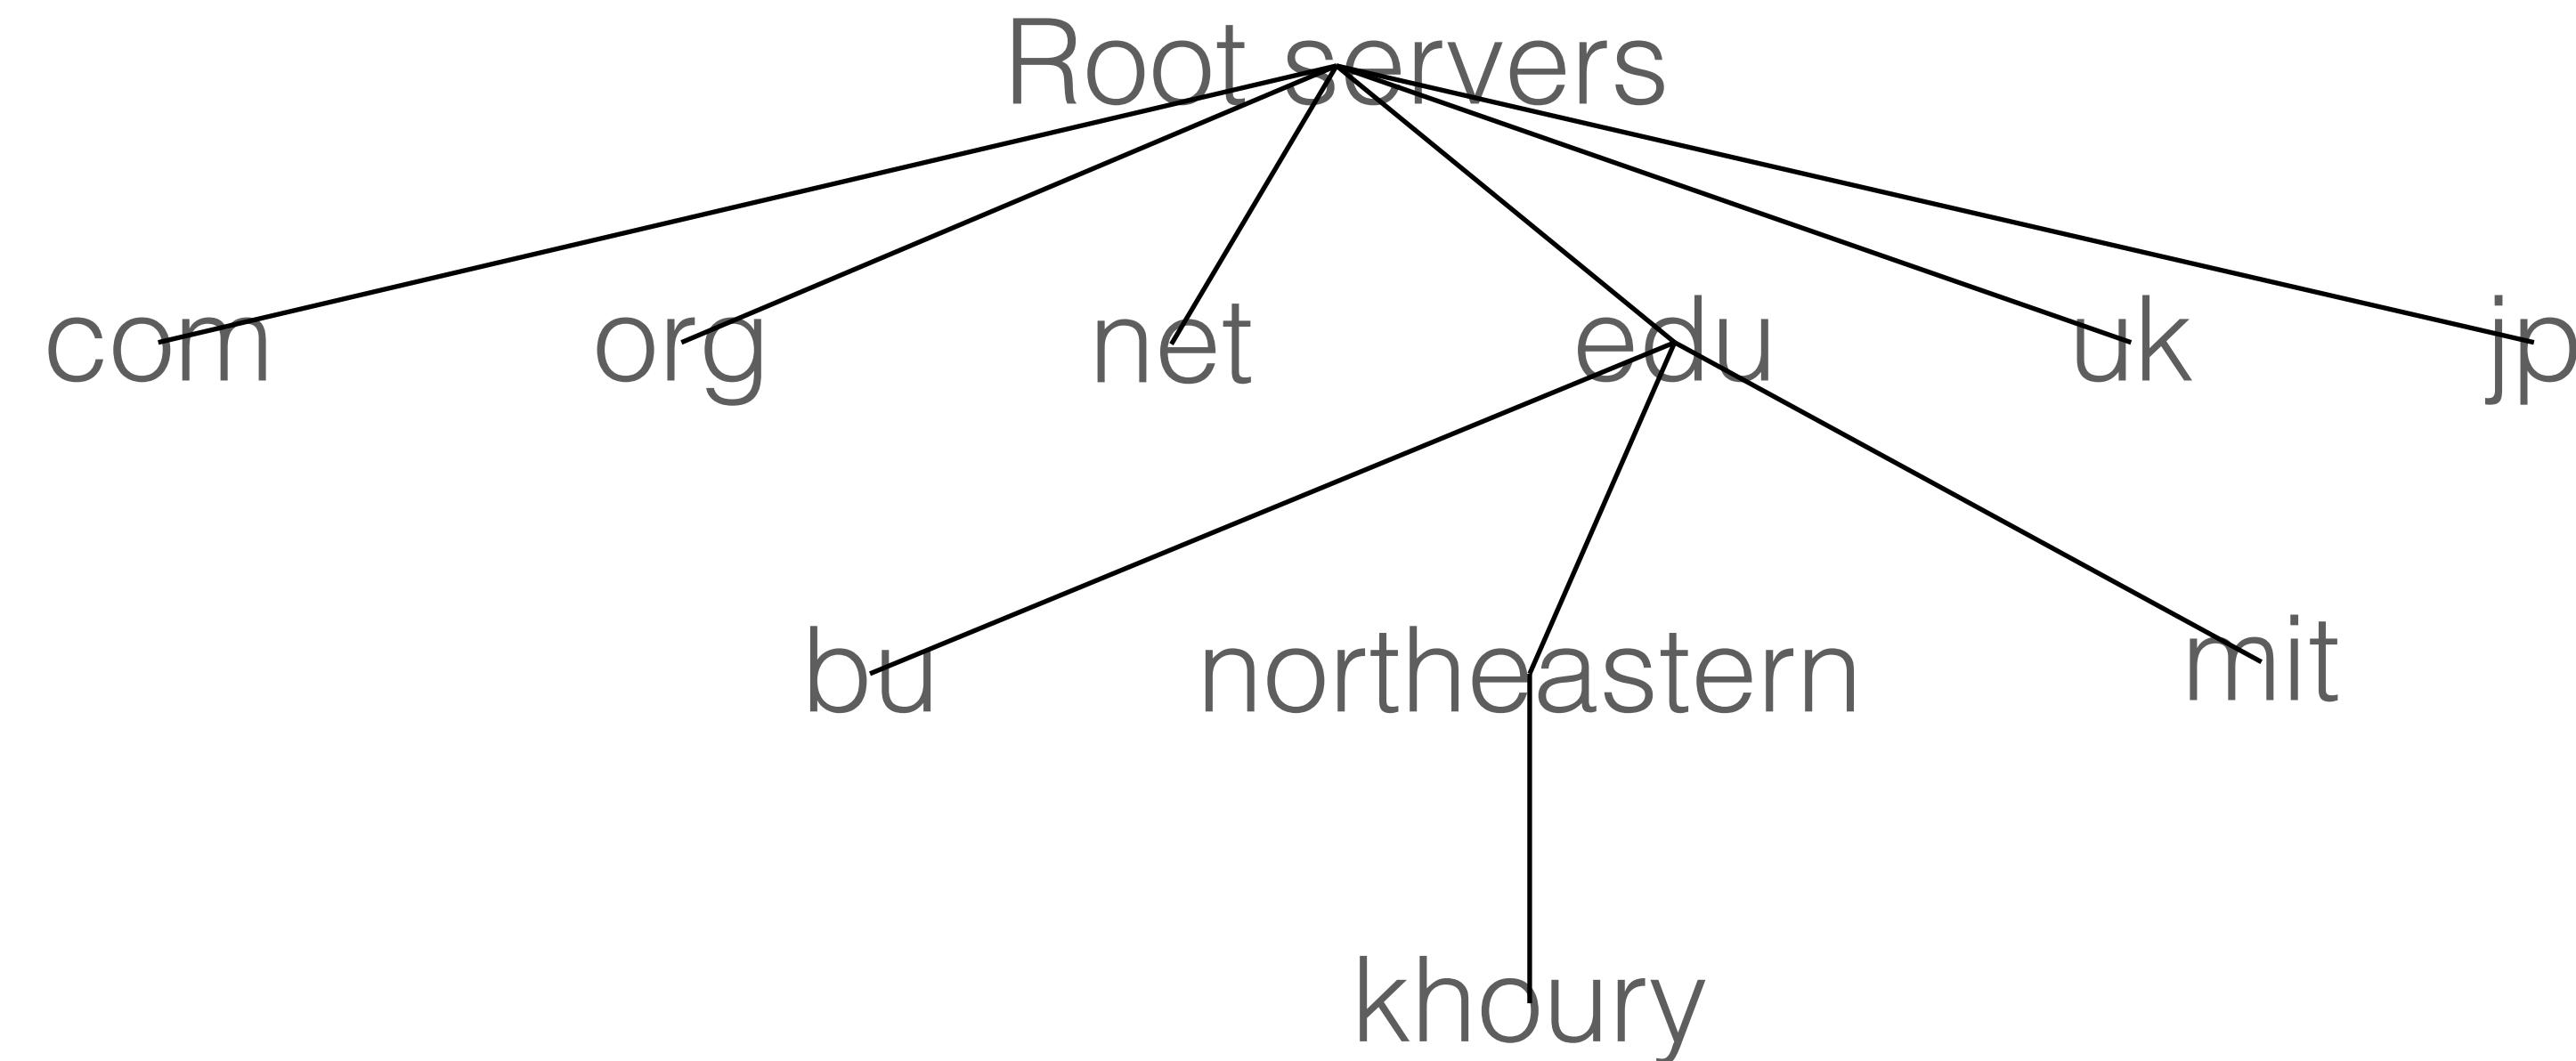
- DNS requirements for querying records:
  - Scalability - grow in number of names
  - Performance - grow in number of requests
  - Latency - provide answers with little delay
  - Fault tolerance & availability - always provide a response
- DNS requirements for updating/creating records:
  - Latency - OK if it takes minutes/hours for an update to take full effect
  - Performance - Expect far fewer writes than reads

# Preliminary Design: How to organize data in the system?

- This depends to a large degree on whether there is shared state
- Usually, there is some shared state
- How important is it to synchronize?
- What about our DNS example?
  - Domains can be split (e.g., .com, .edu, .info, .eu, .jp, ...)
  - Huge volume of requests – multiple nodes need to provide the same mappings, consistently

# How to organize DNS

Idea: break apart responsibility for each part of a domain name (**zone**) to a different group of servers



Each zone is a continuous section of name space  
Each zone has an associate set of name servers

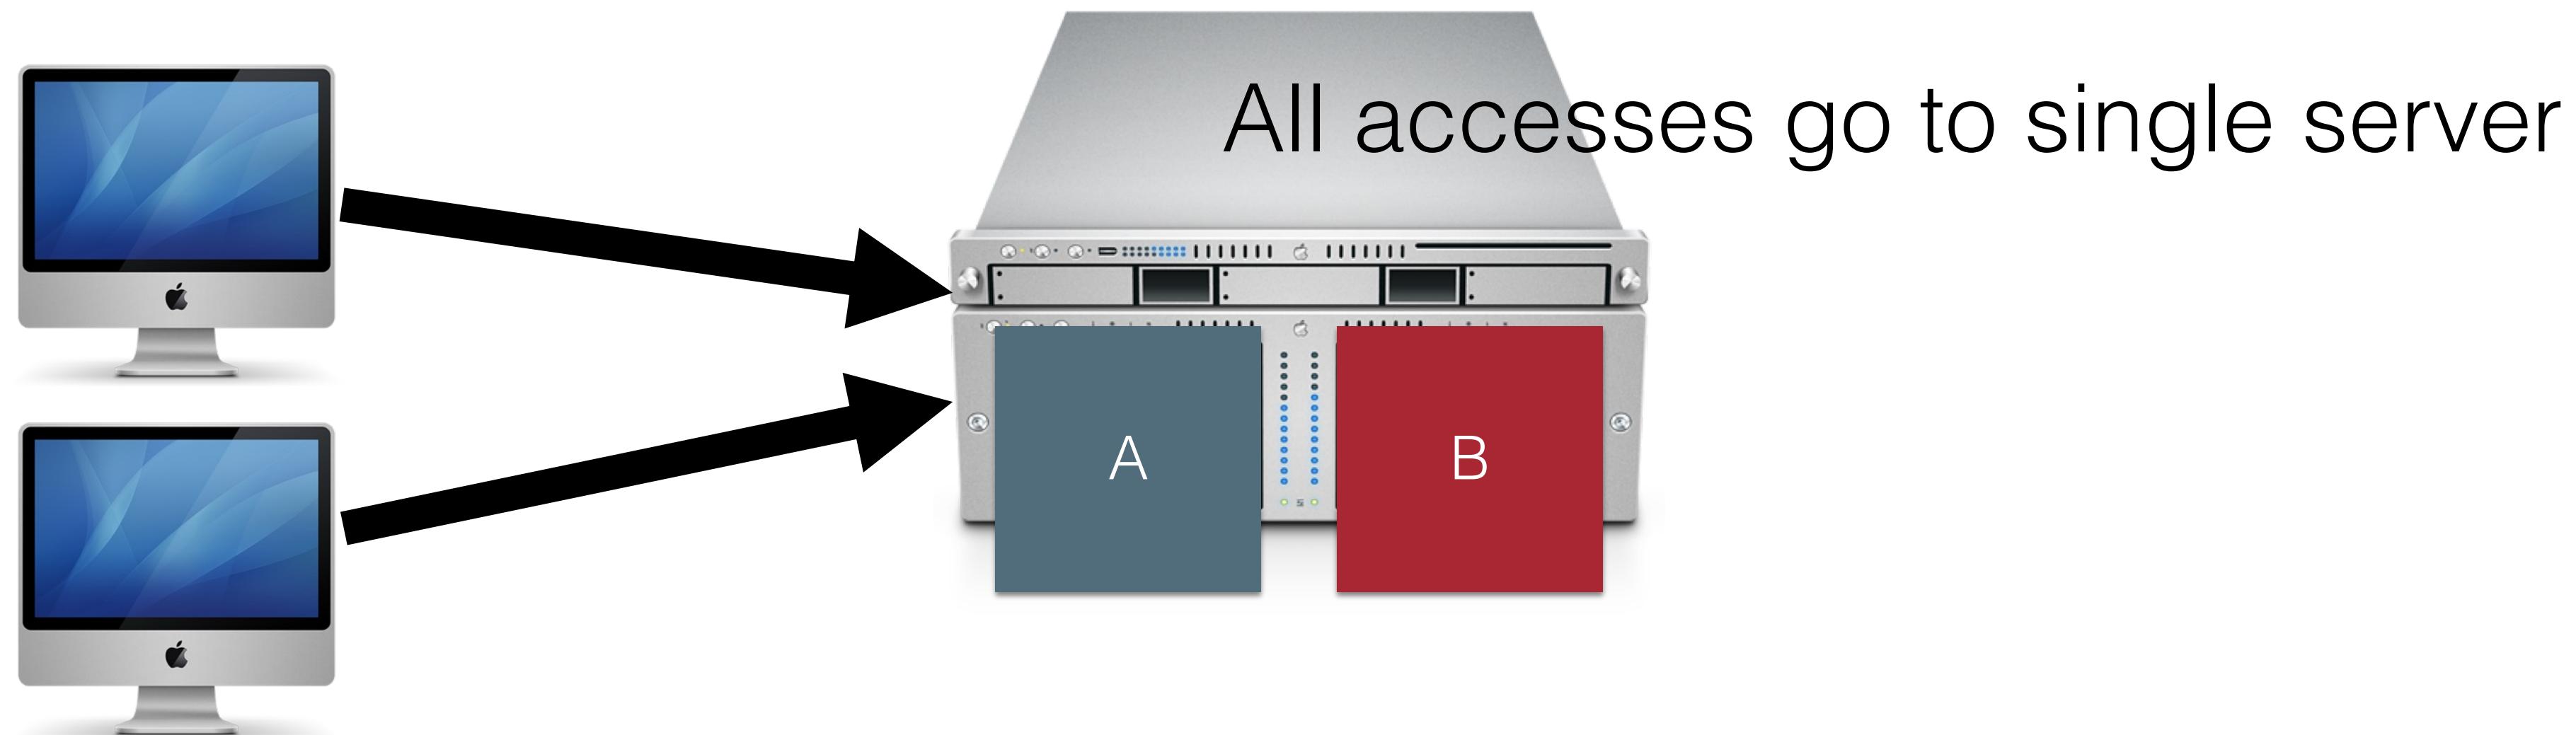
# How to organize DNS

Idea: break apart responsibility for each part of a domain name (**zone**) to a different group of servers

In other words, we **partition** the domain names according to the top-level domain.

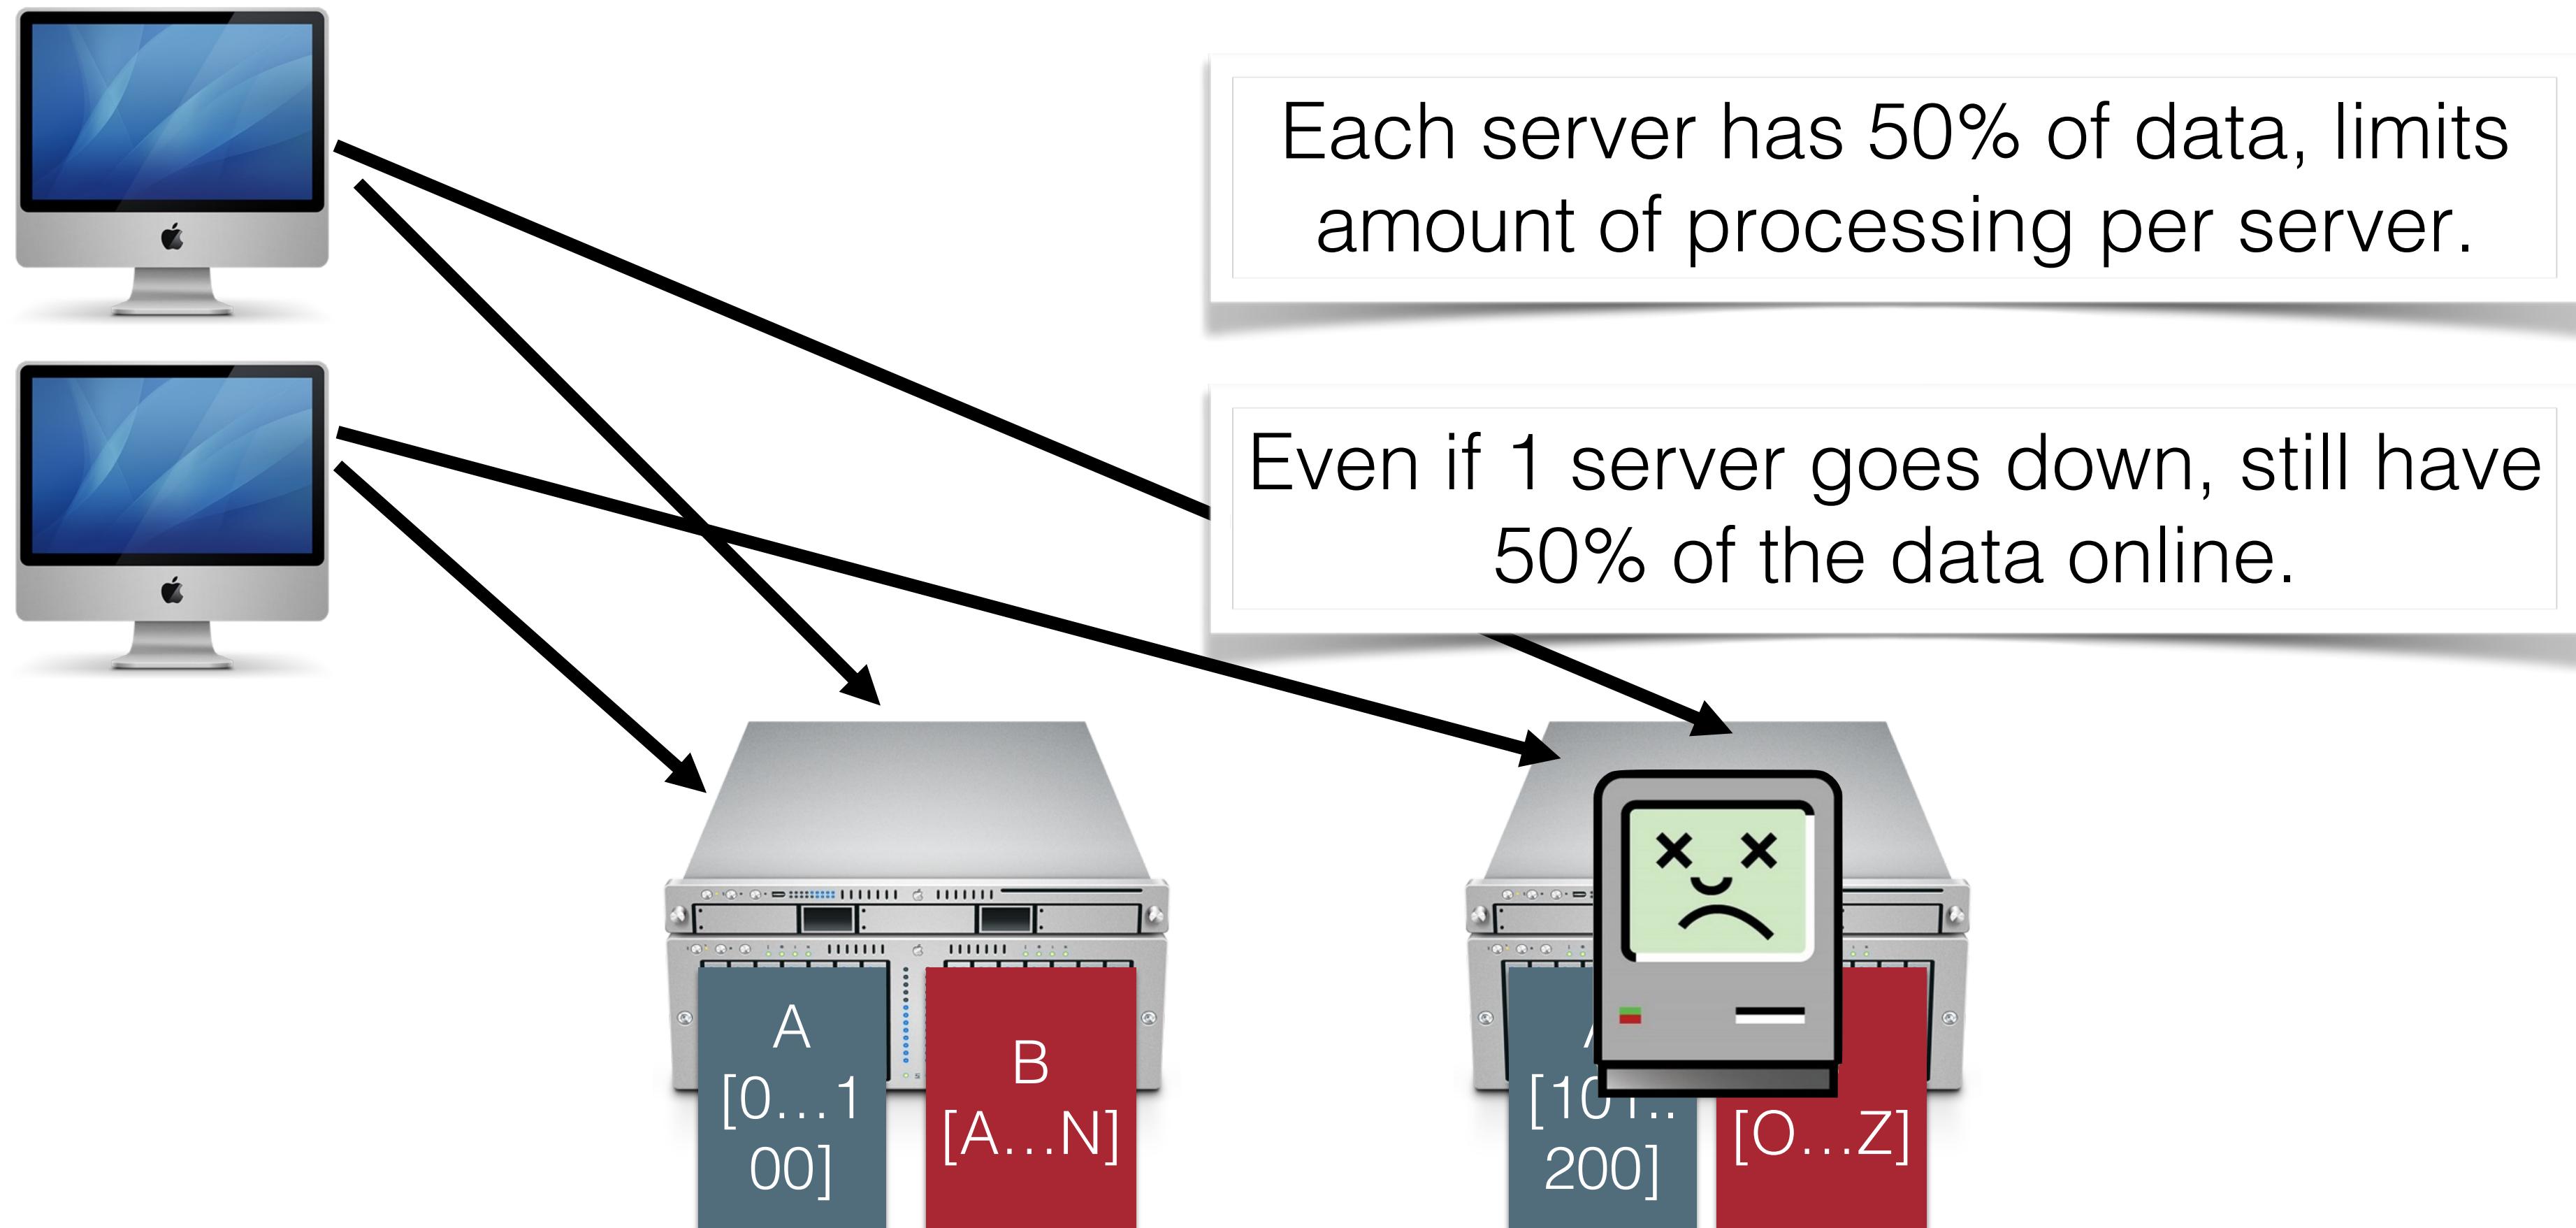
# Recurring Solution #1: Partitioning

- Partitioning is a common strategy to distributing a system and its data
- Starting from a non-distributed system:

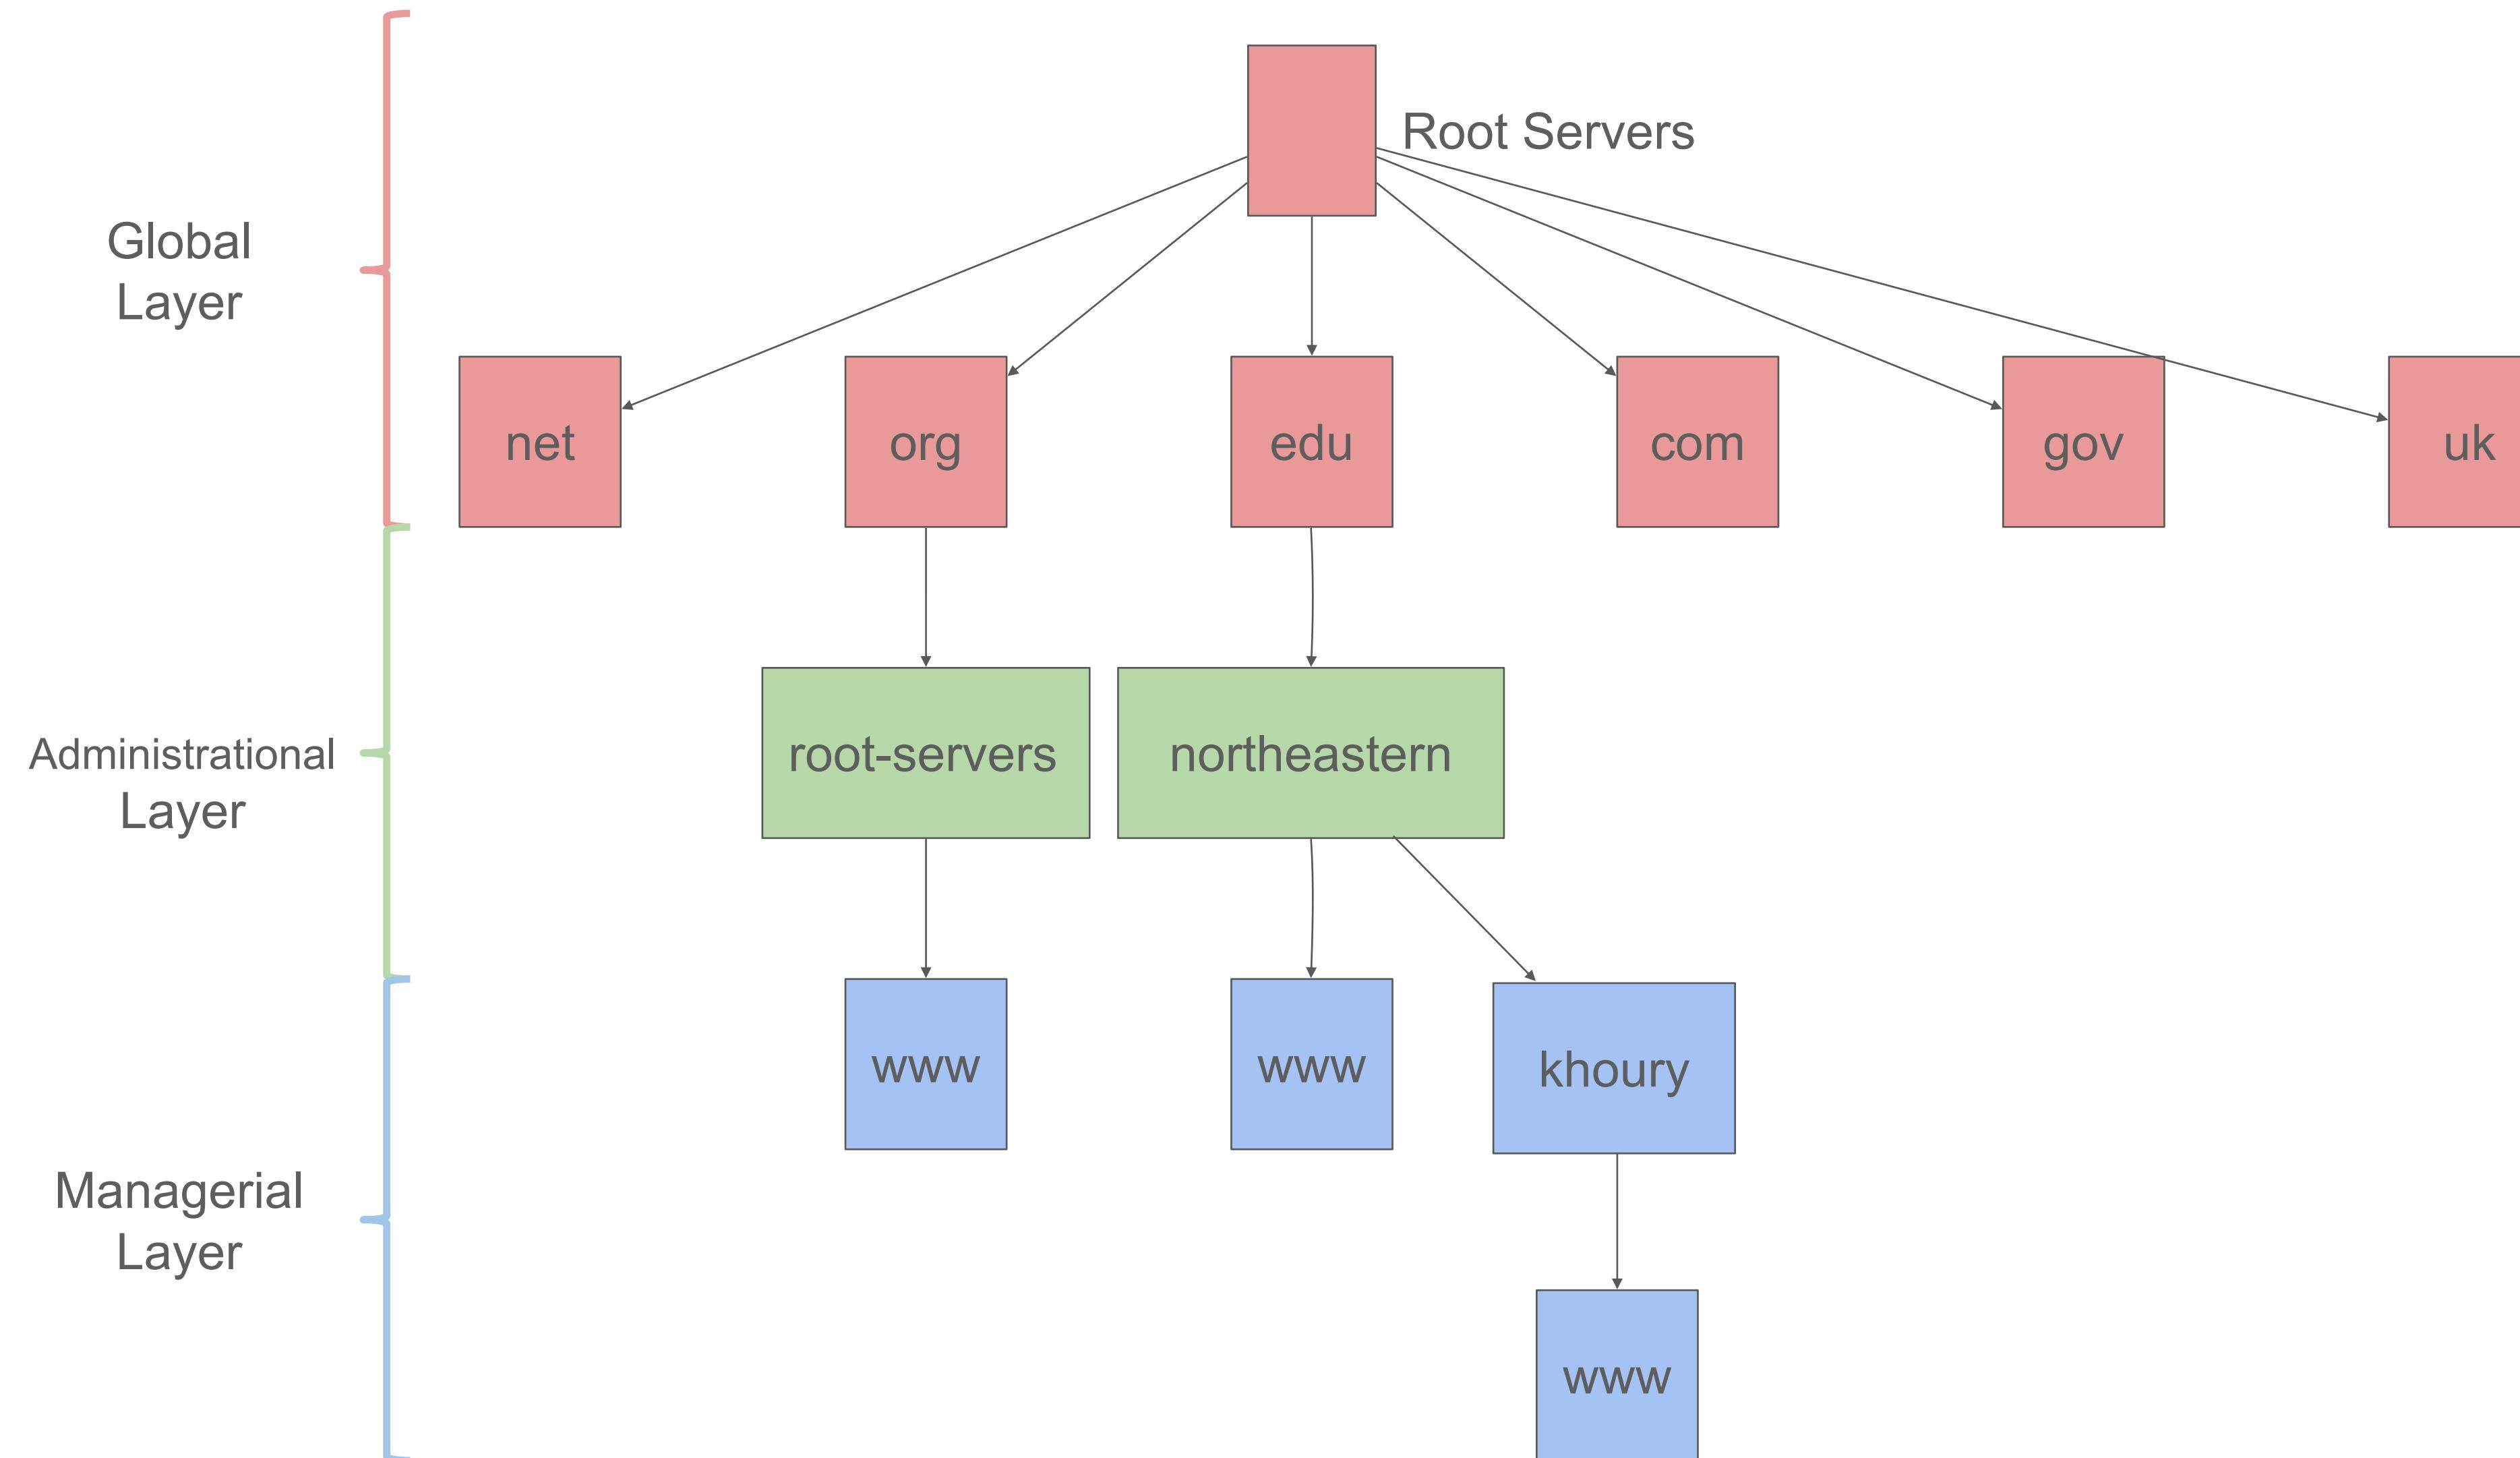


# Recurring Solution #1: Partitioning

- Divide data up in some (hopefully logical) way
- Makes it easier to process data concurrently (cheaper reads)

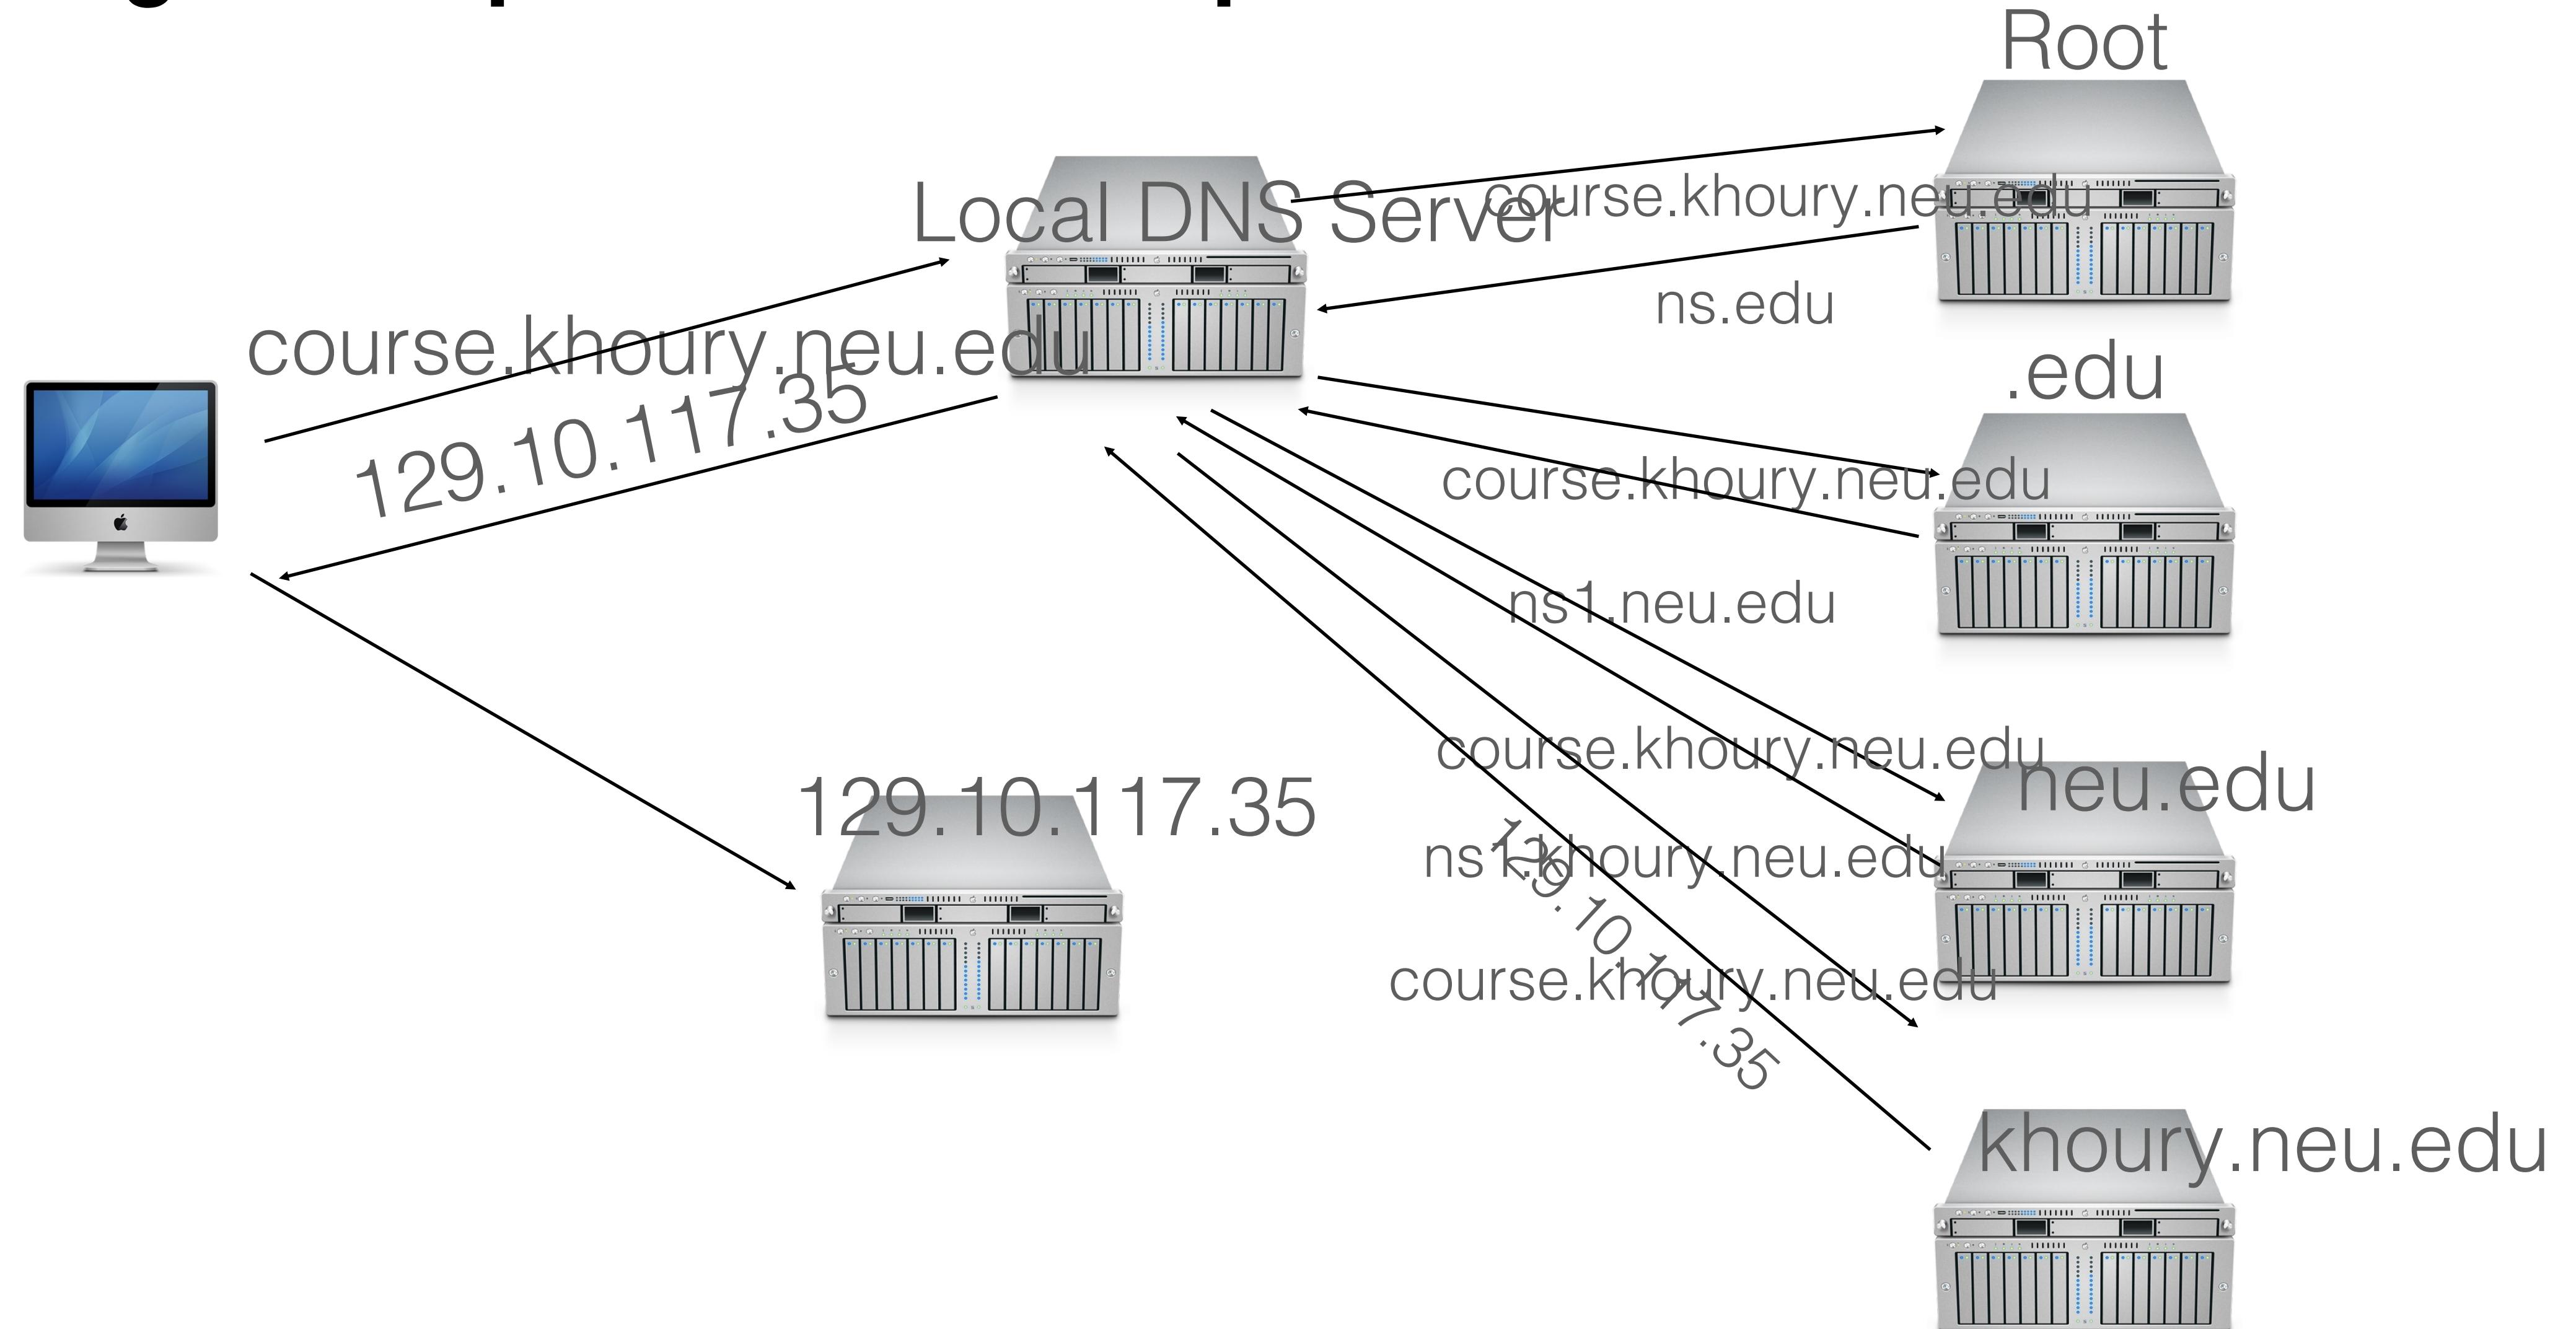


# Partitioning DNS



# DNS: Example

What might a request look like in practice?

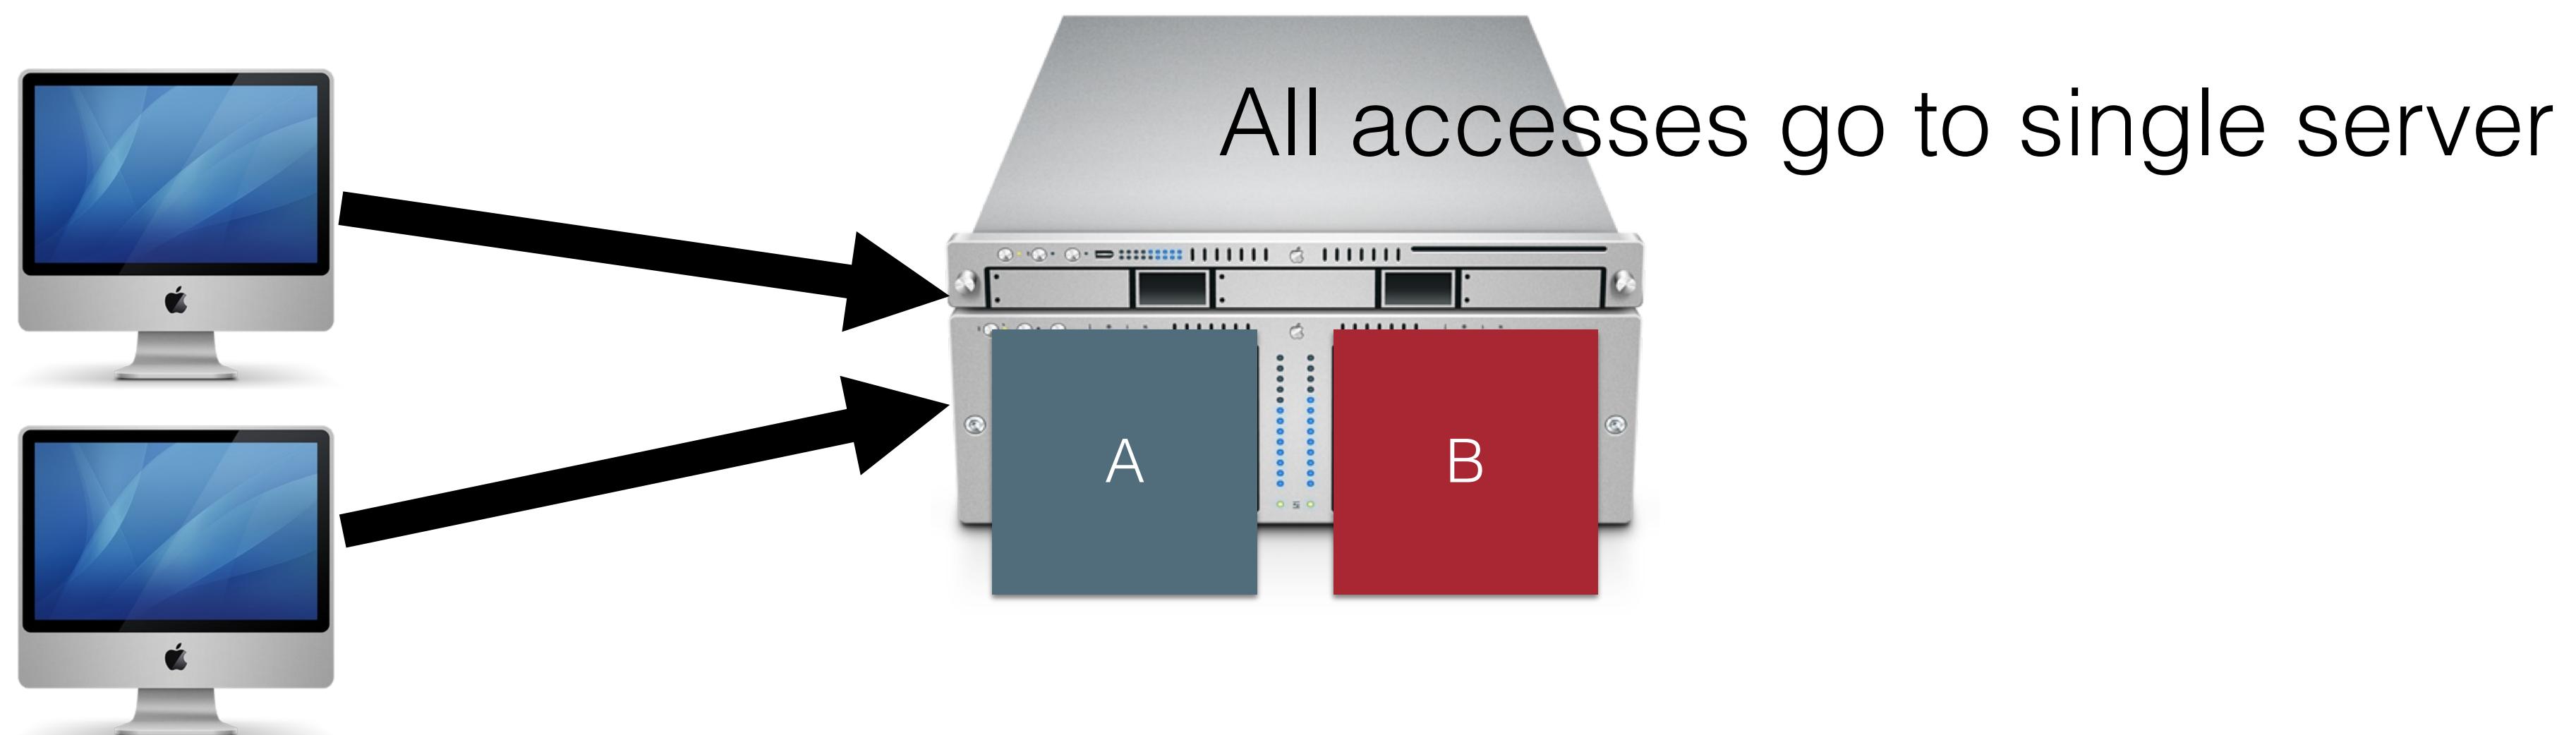


# How to deal with volume?

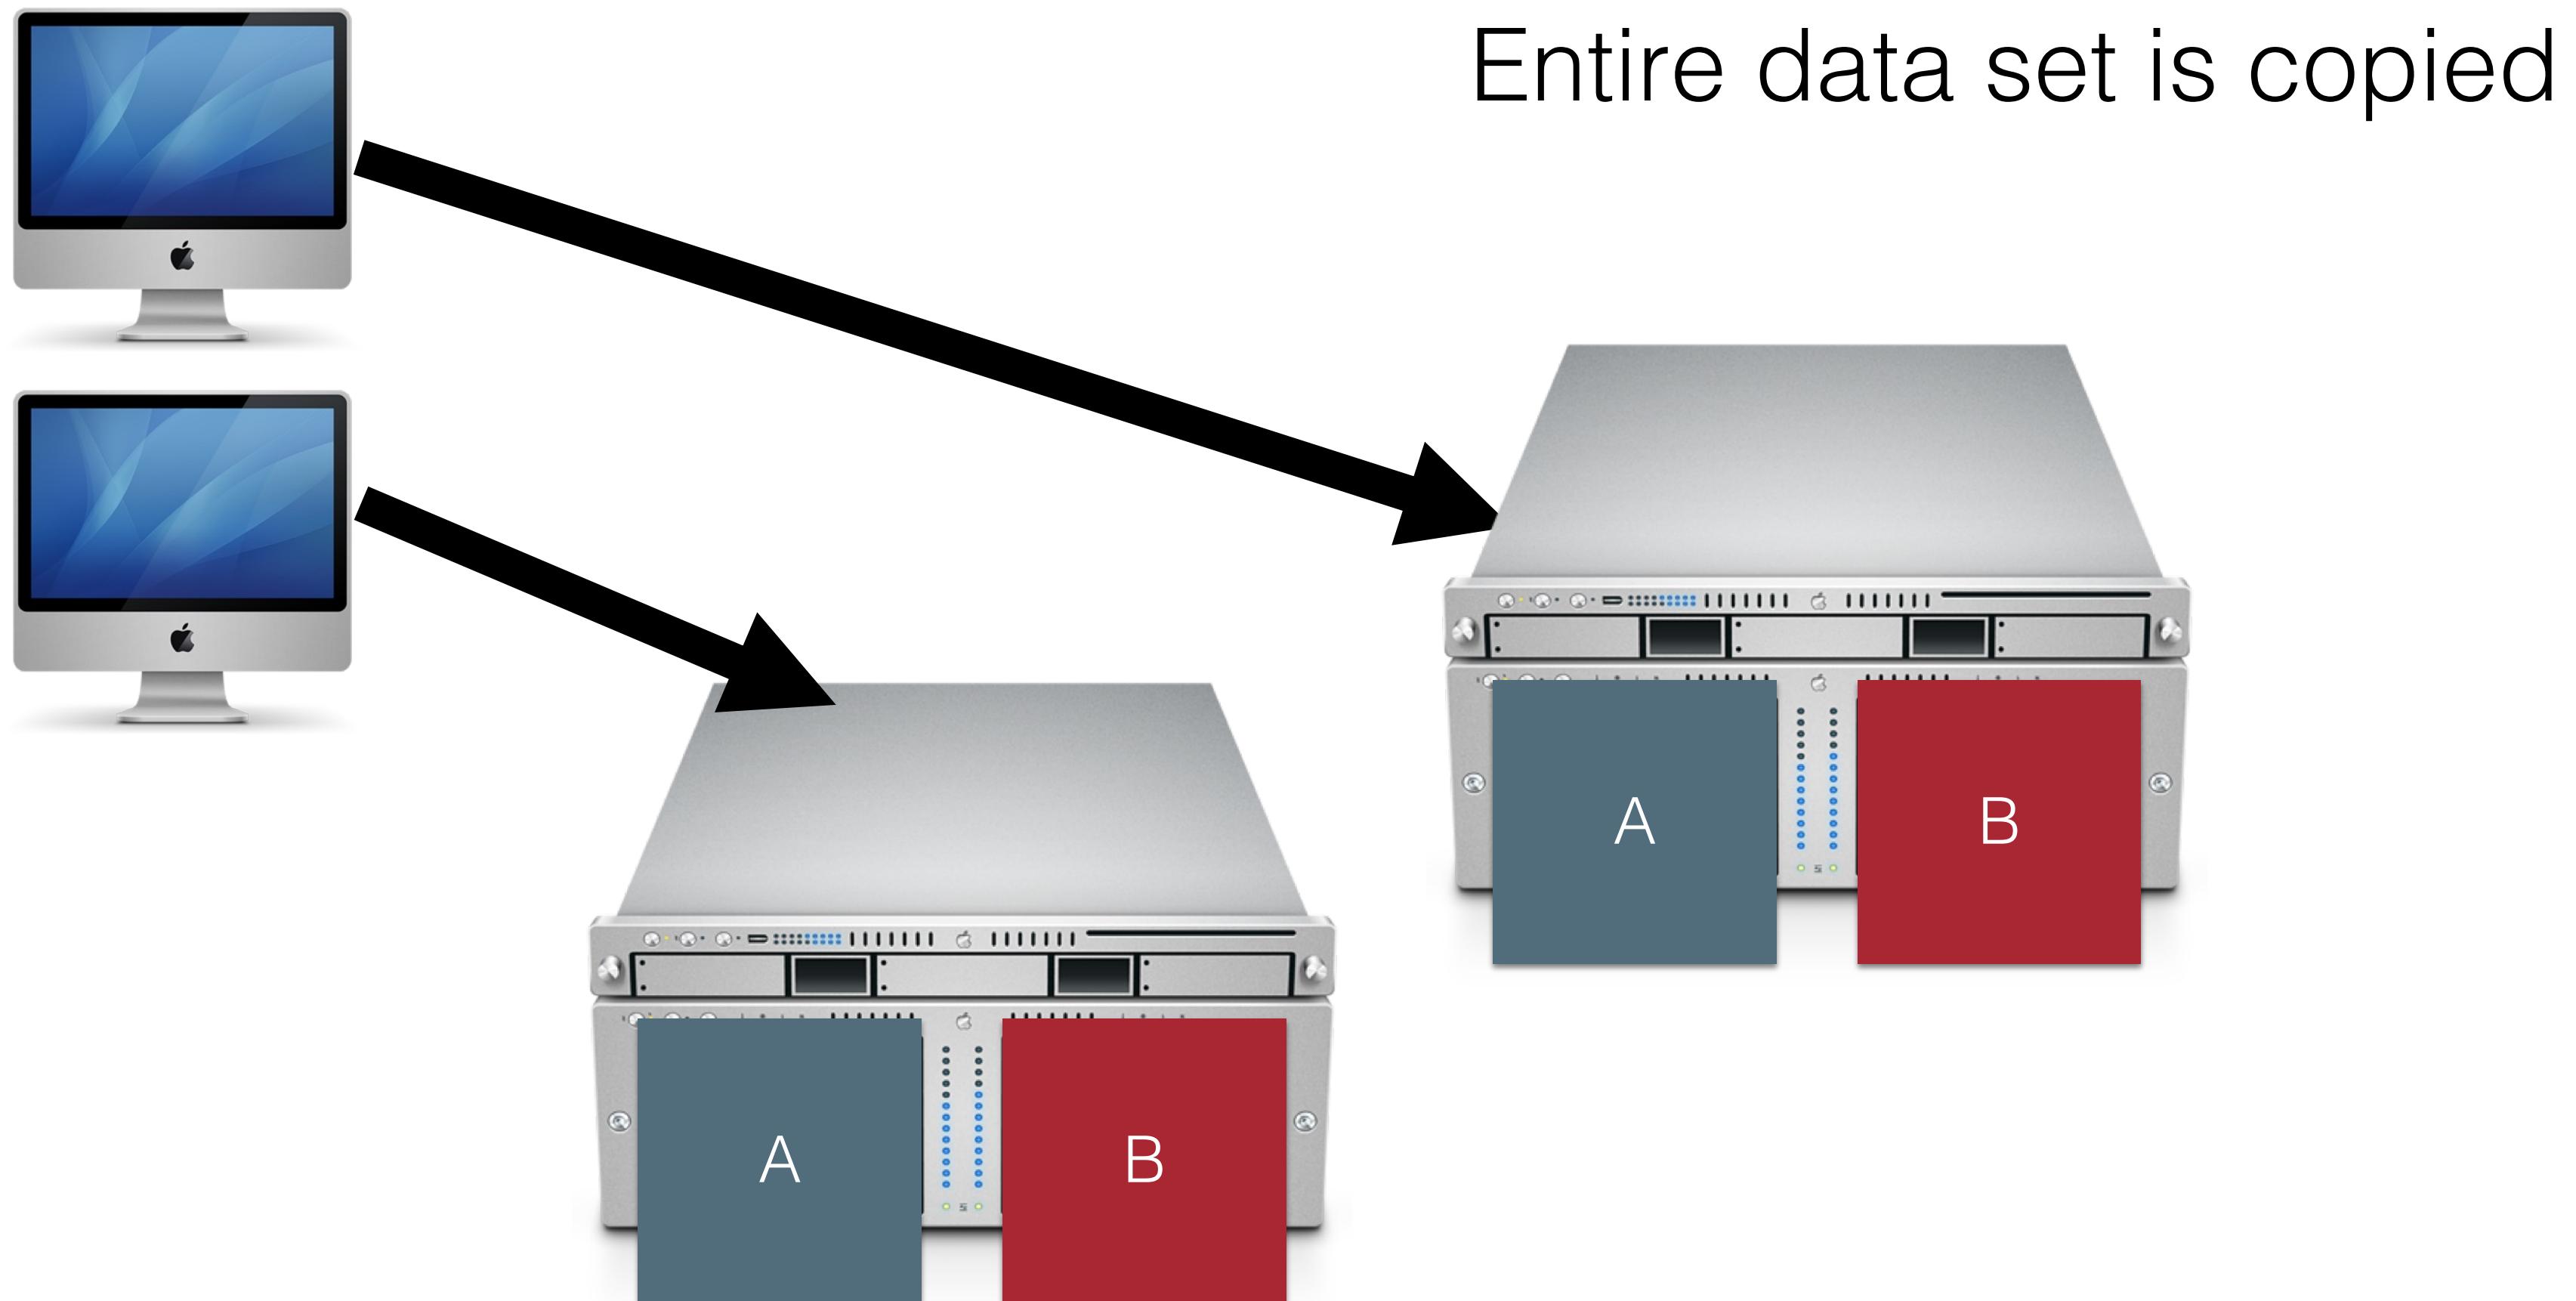
- We successfully distributed requests following the hierarchical nature of domain names
- However, e.g., .com is a very popular TLD – there might be (hundreds of) thousands of requests happening at any given time
- We may need several nodes just servicing .com
- This leads to **replication**

# Recurring Solution #2: Replication

- Goal: Any node should be able to process any request
- Again, starting from a non-distributed system:



# Recurring Solution #2: Replication



# Recurring Solution #2: Replication

- Improves performance:
  - Client load can be evenly shared between servers
  - Reduces latency: can place copies of data nearer to clients
- Improves availability:
  - One replica fails, still can serve all requests from other replicas

# Replication in DNS – Root Servers

- 13 root servers
  - [a-m].root-servers.org
  - E.g., d.root-servers.org
- Handled by 12 distinct entities
  - (“a” and “j”) are both Verisign
    - Don’t ask why.

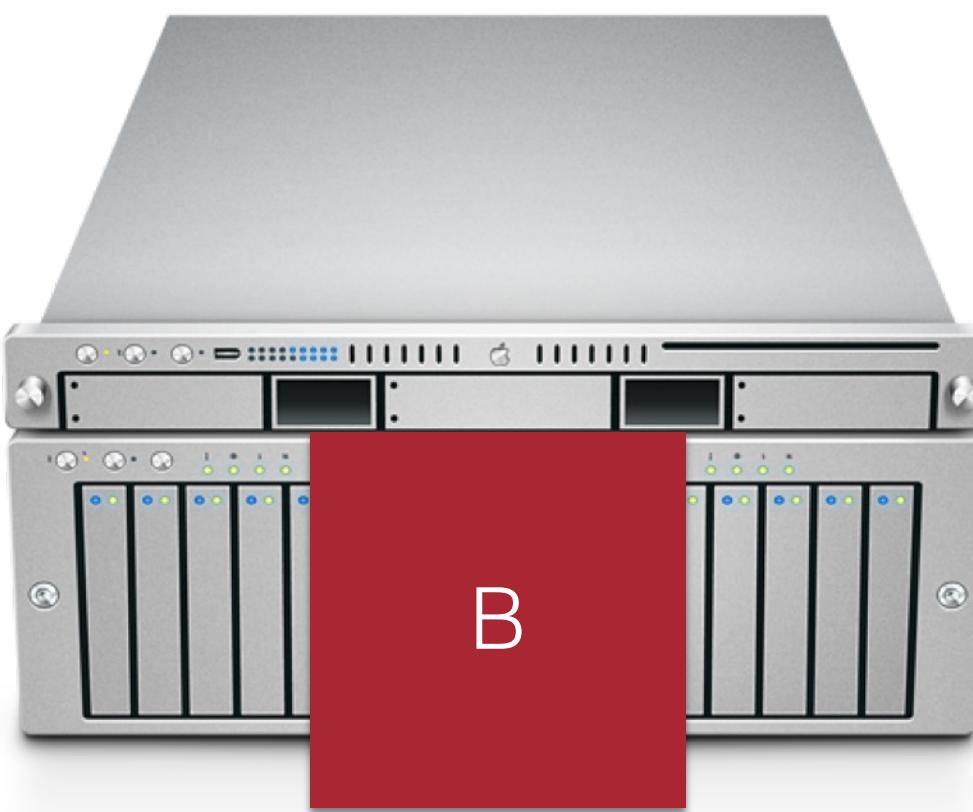
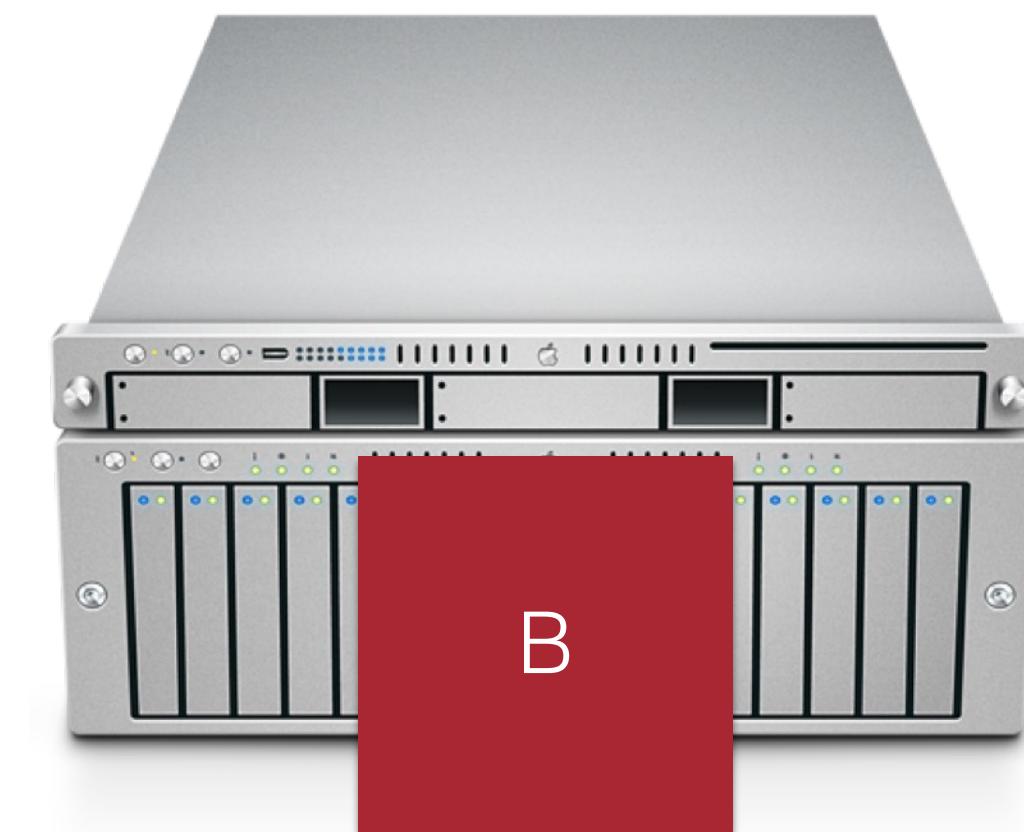
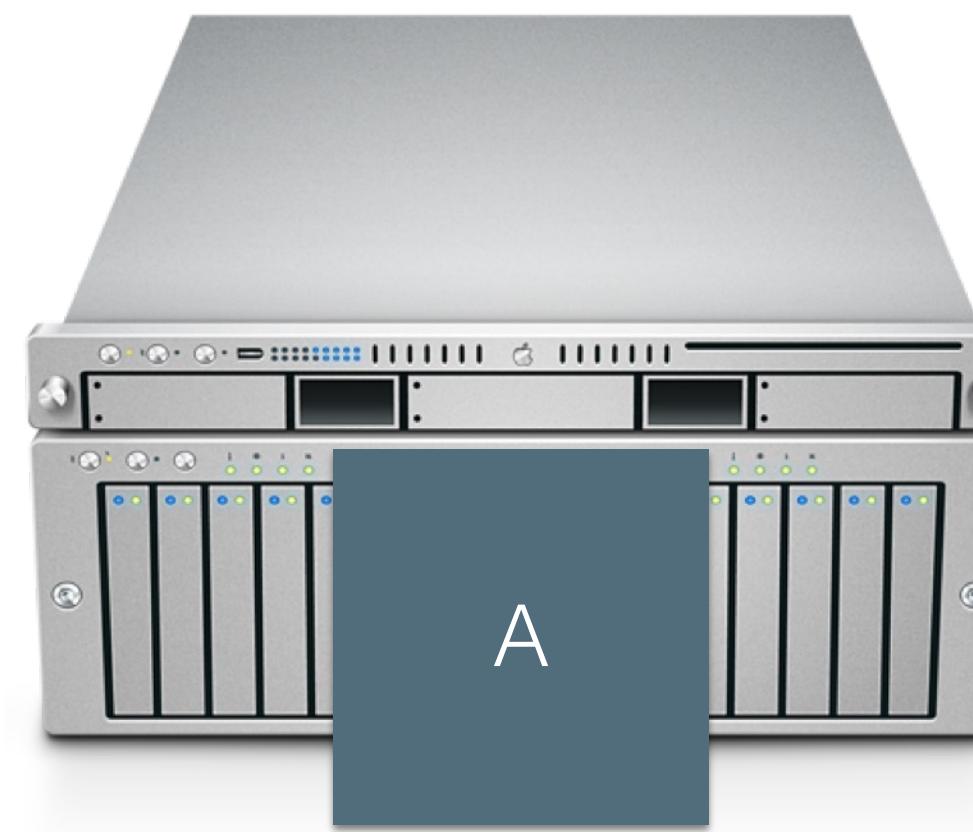
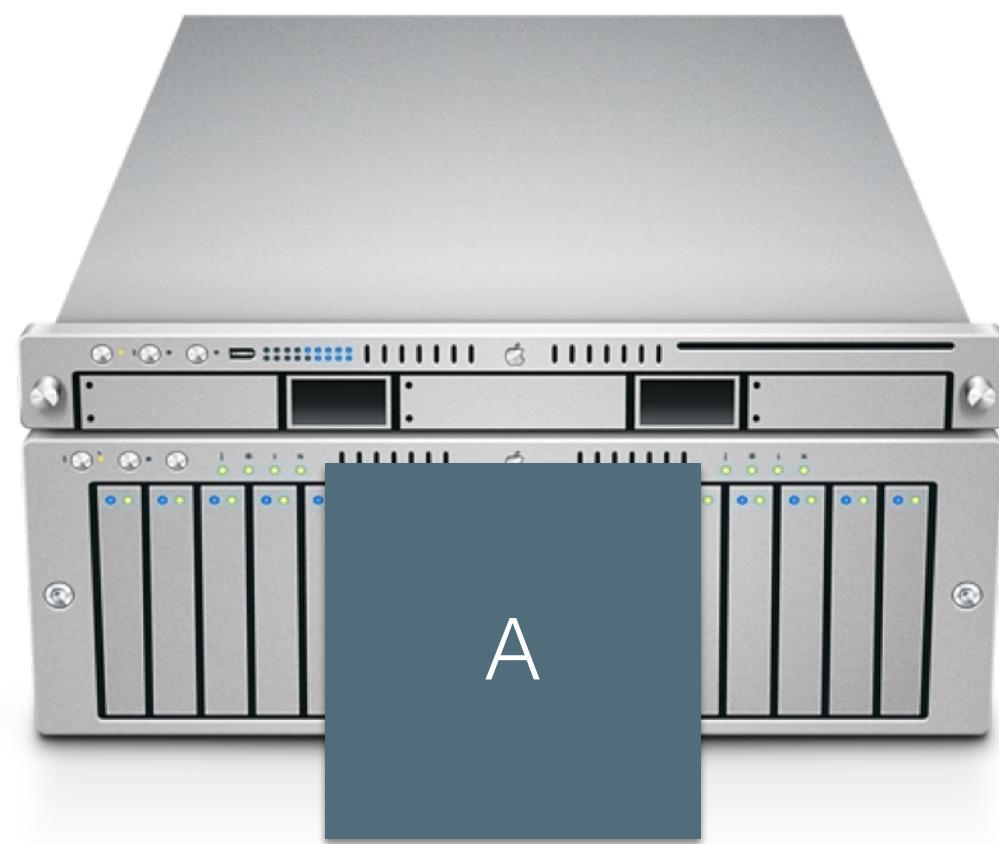
Verisign, Inc.	a
Information Sciences Institute	b
Cogent Communications	c
University of Maryland	d
NASA Ames Research Center	e
Internet Systems Consortium, Inc.	f
U.S. DOD Network Information Center	g
U.S. Army Research Lab	h
Netnod	i
Verisign, Inc.	j
RIPE NCC	k
ICANN	l
WIDE Project	m

# There is replication even within the root servers

- 13 root servers
  - [a-m].root-servers.org
  - E.g., d.root-servers.org
- But each root server has multiple copies of the database, which need to be kept in sync.
- Somewhere around 1500 replicas in total.

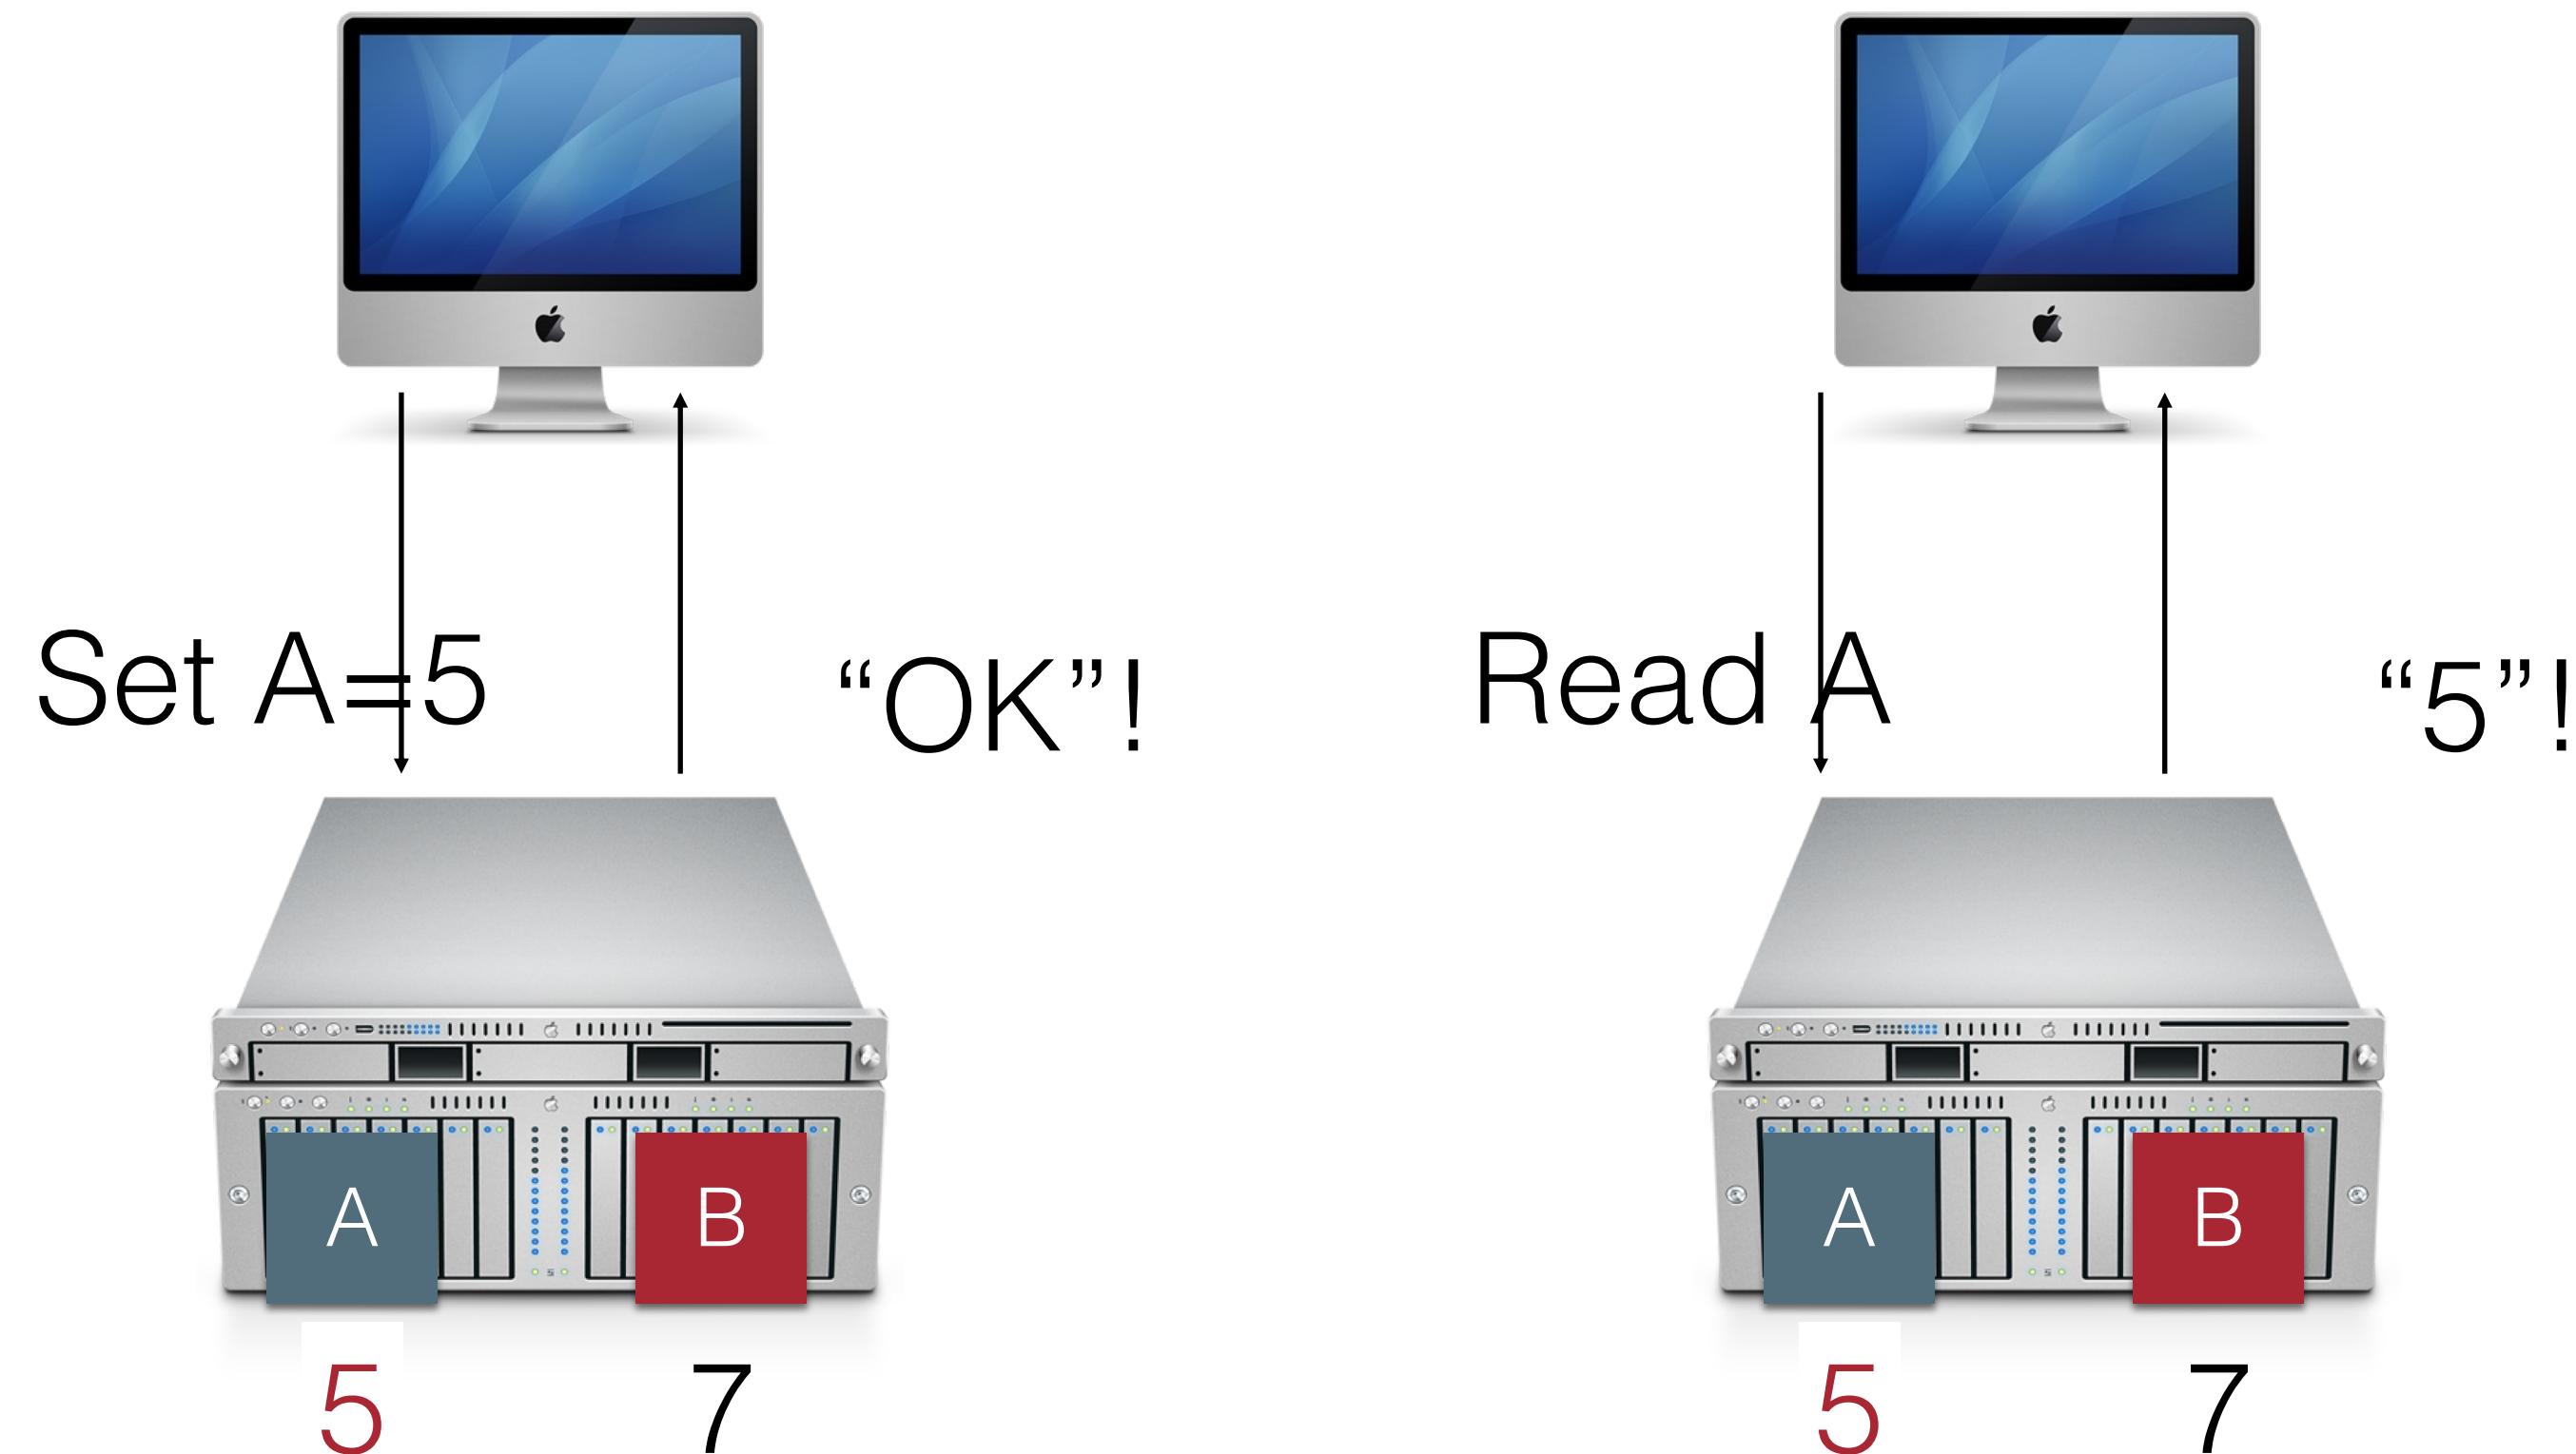
# Partitioning + Replication

- So, DNS combines both partitioning and replication
- As do most distributed systems



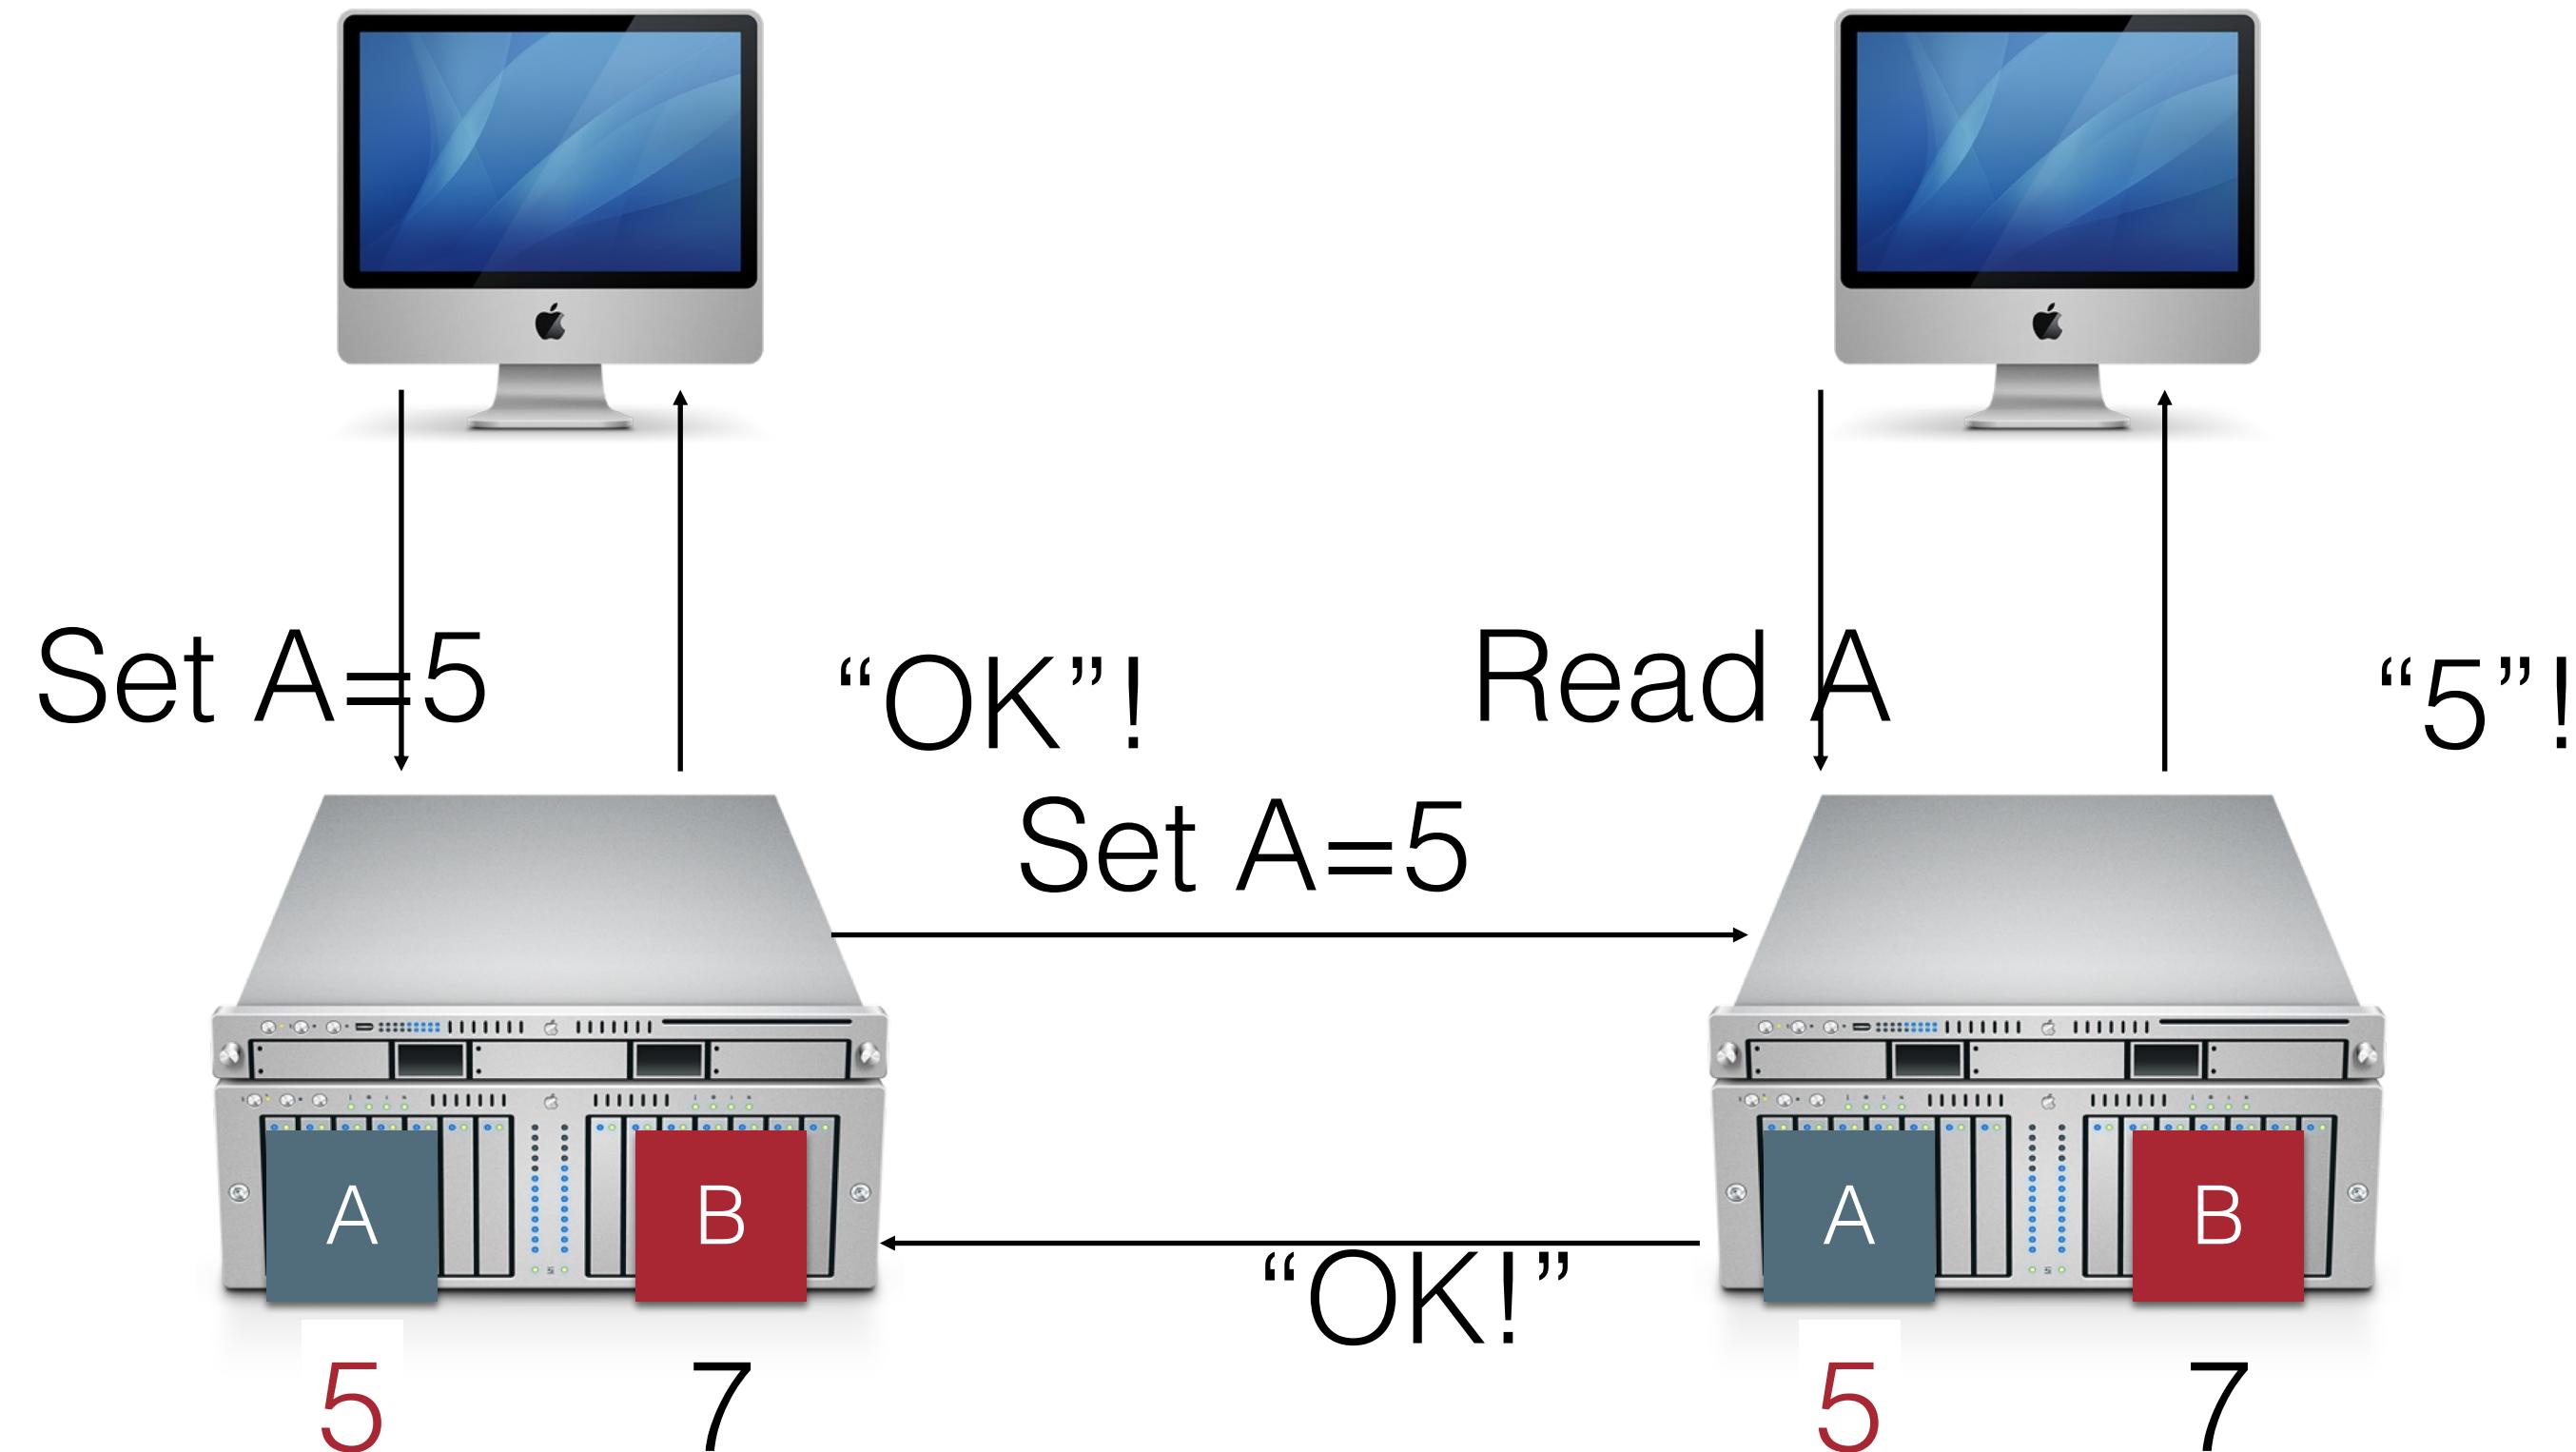
# Replication Problem: Consistency

We probably want our system to work like this



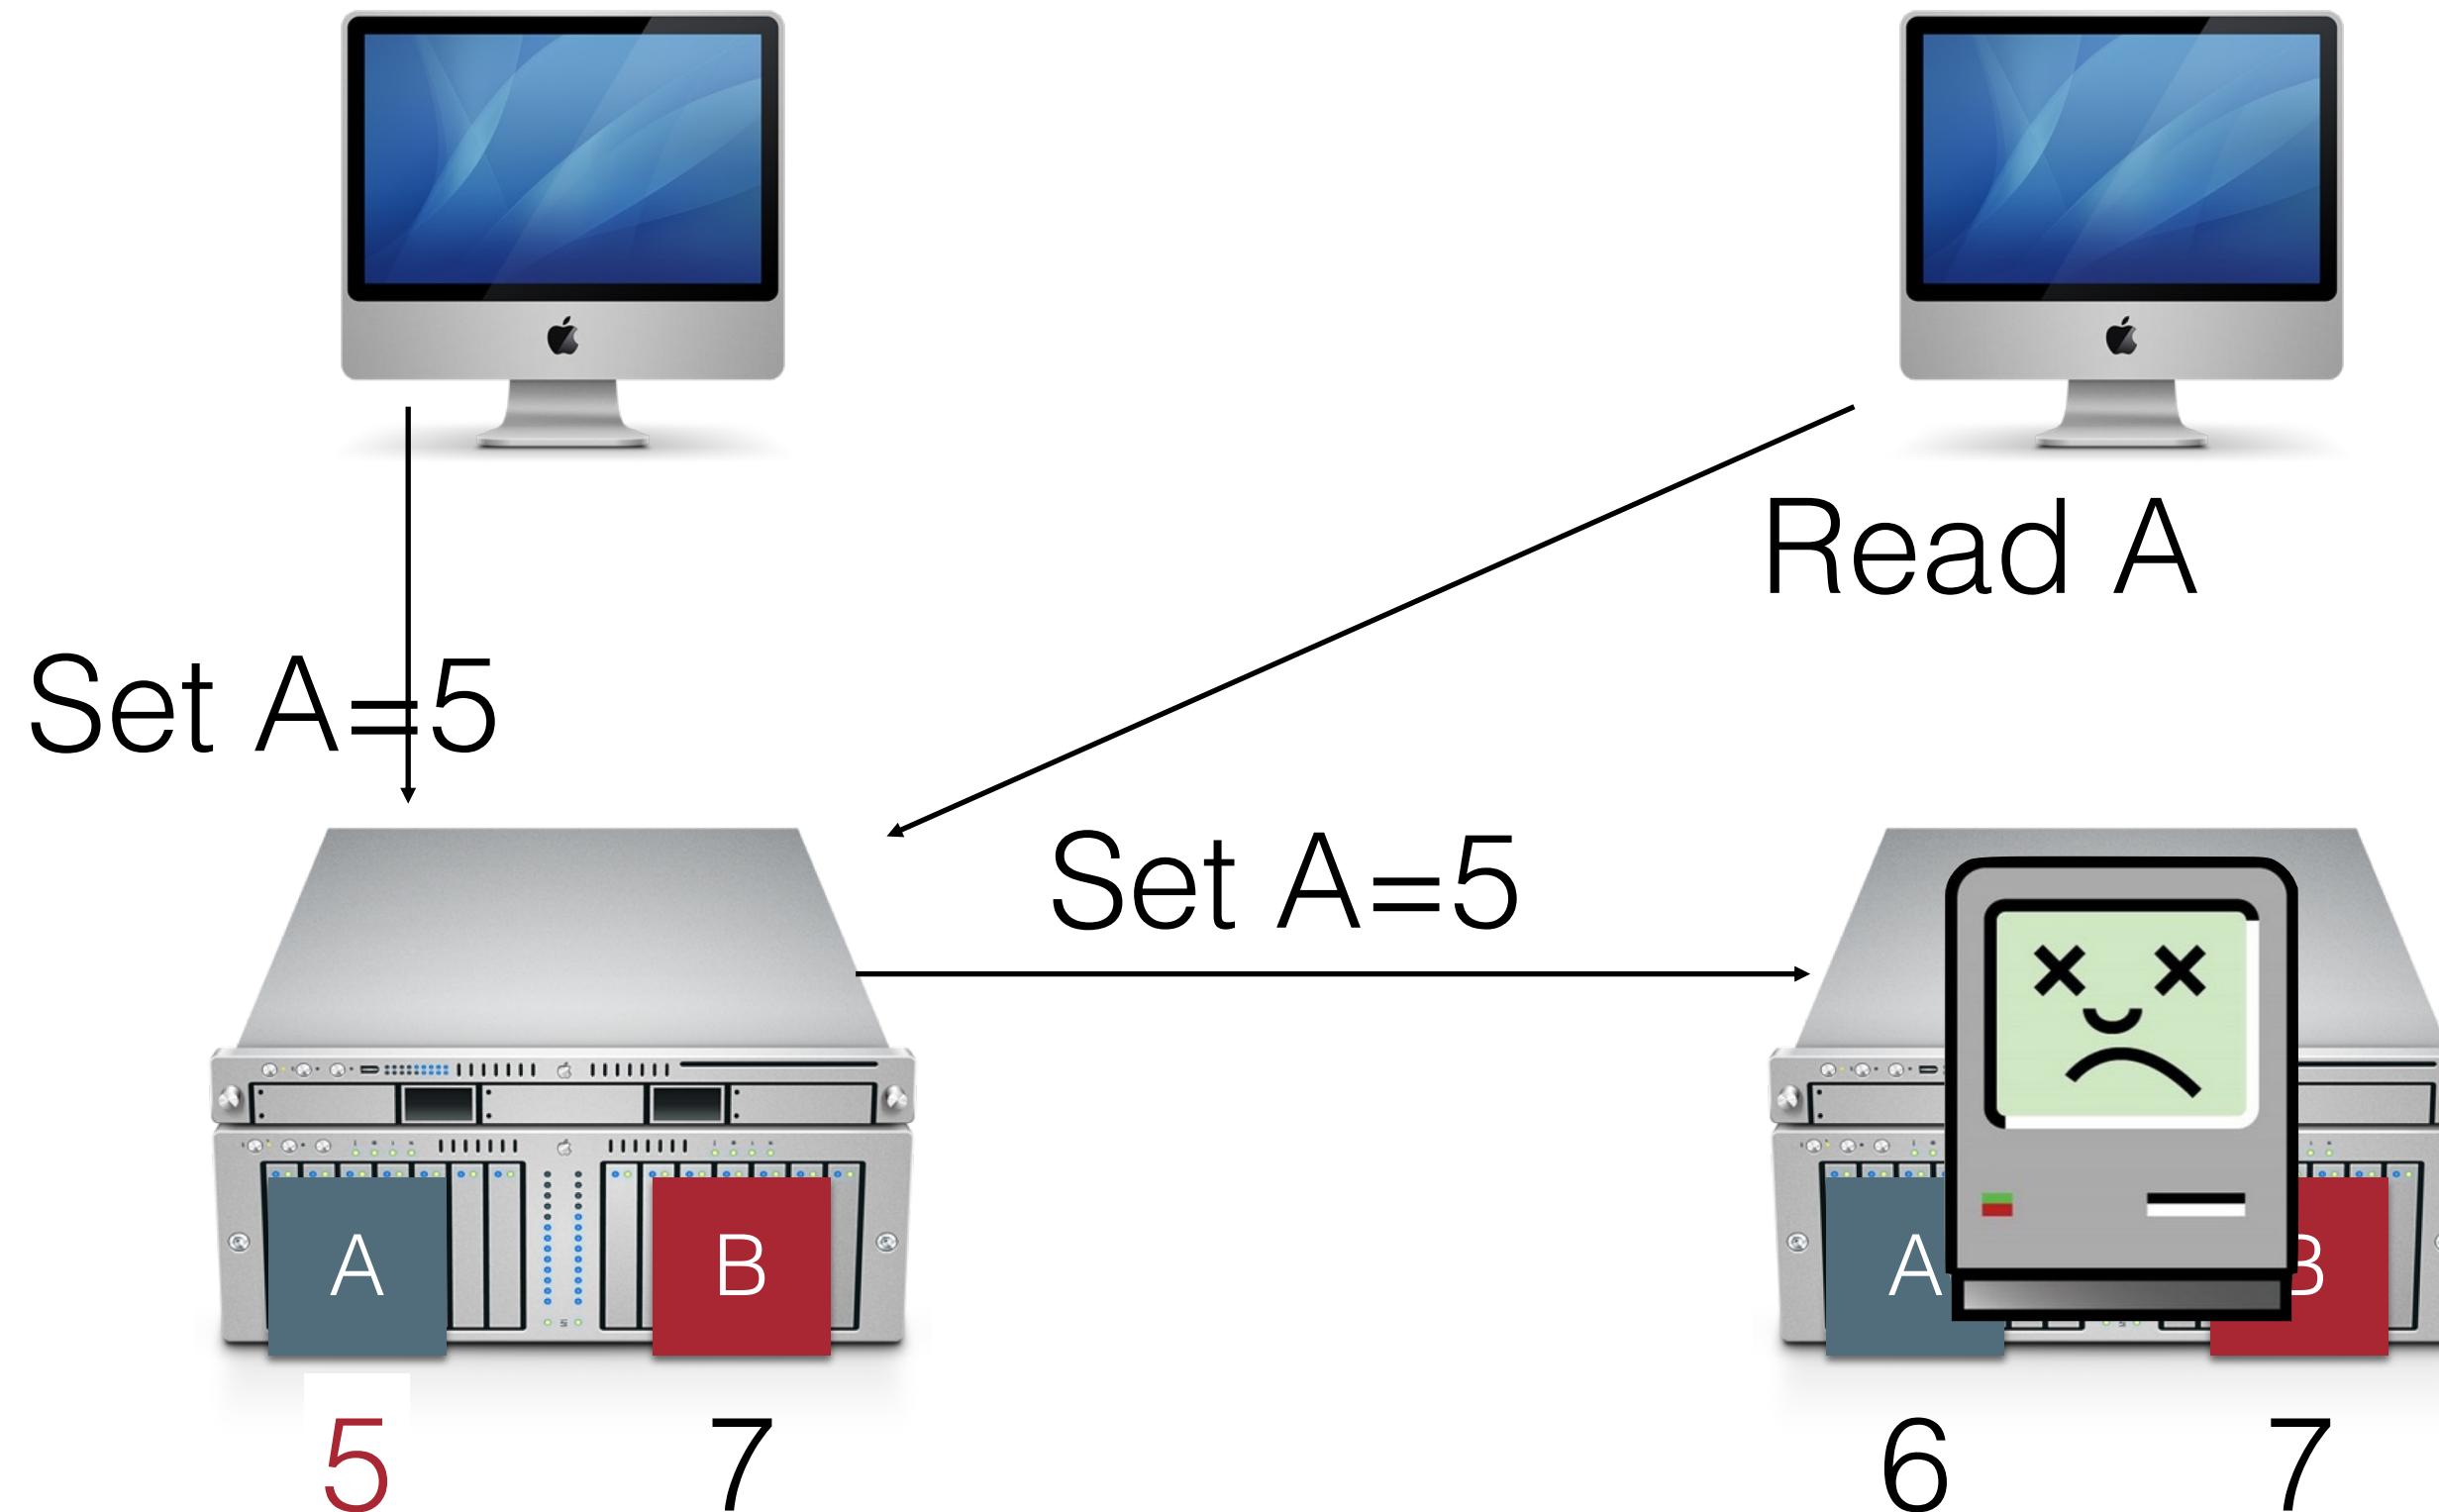
# Sequential Consistency

AKA: Behaves like a single machine would



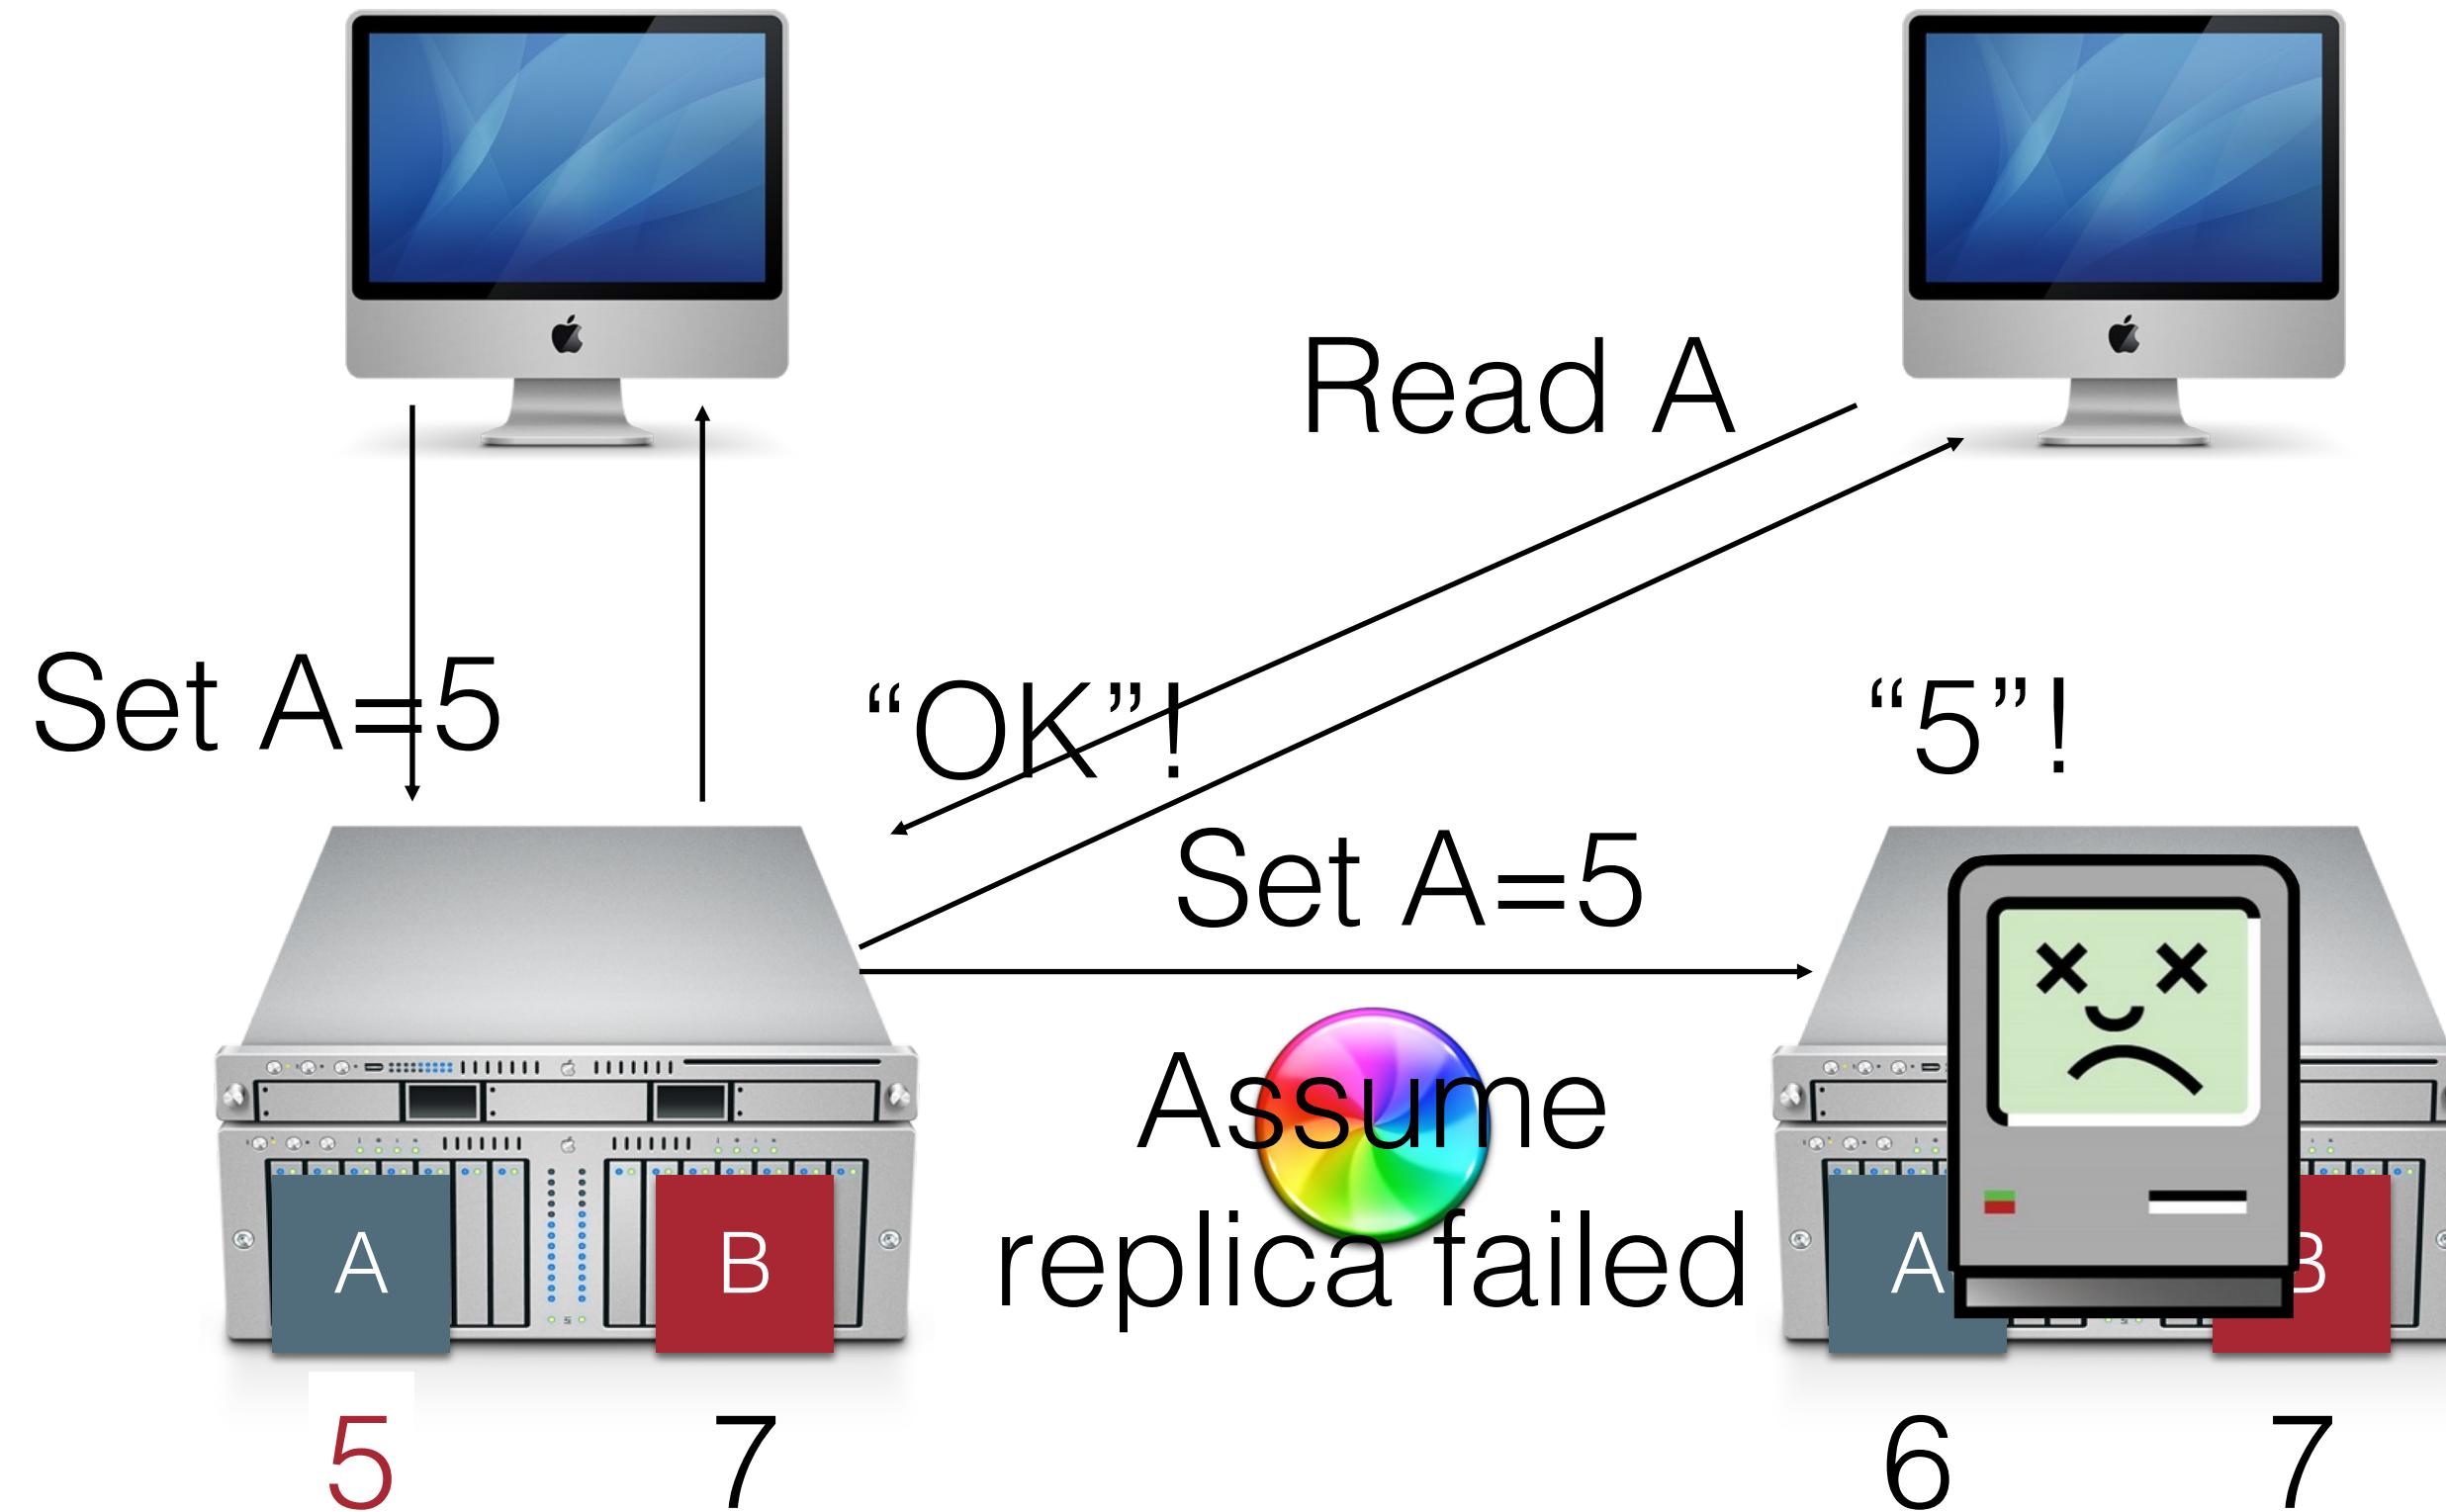
# Availability

If at least one node is online, can we still answer a request?

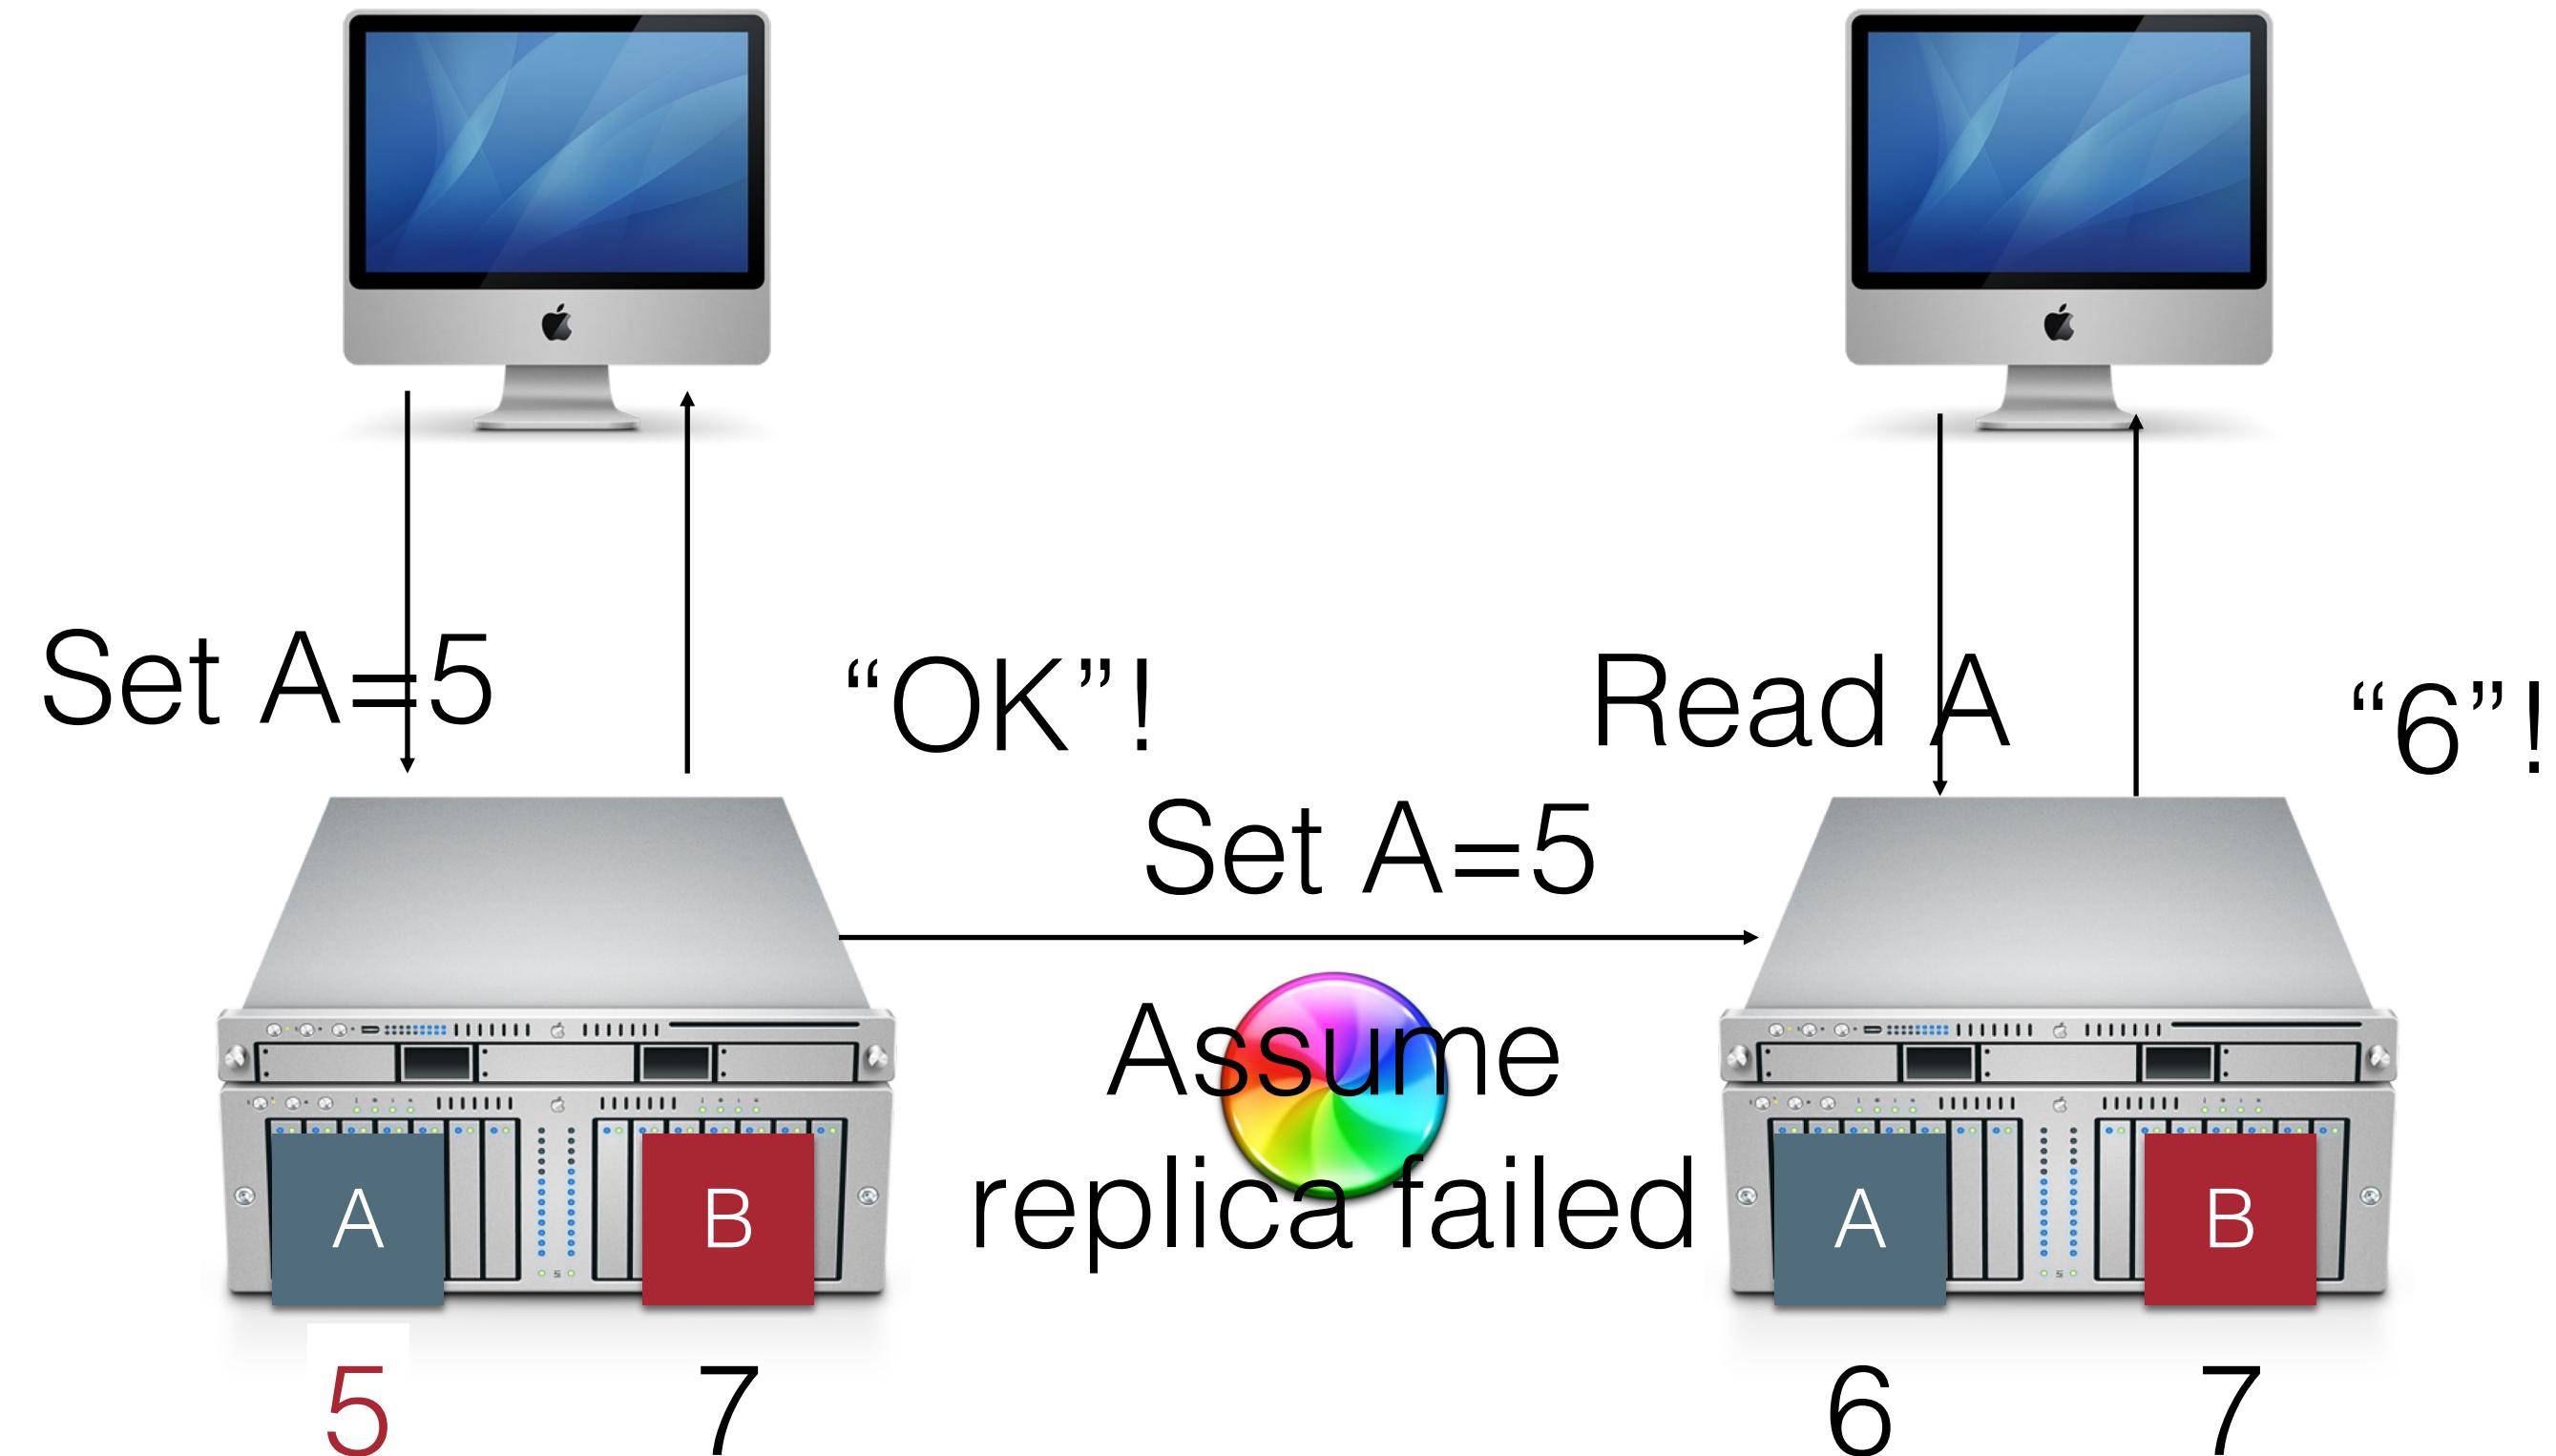


# Consistent + Available

On timeout, assume node is crashed



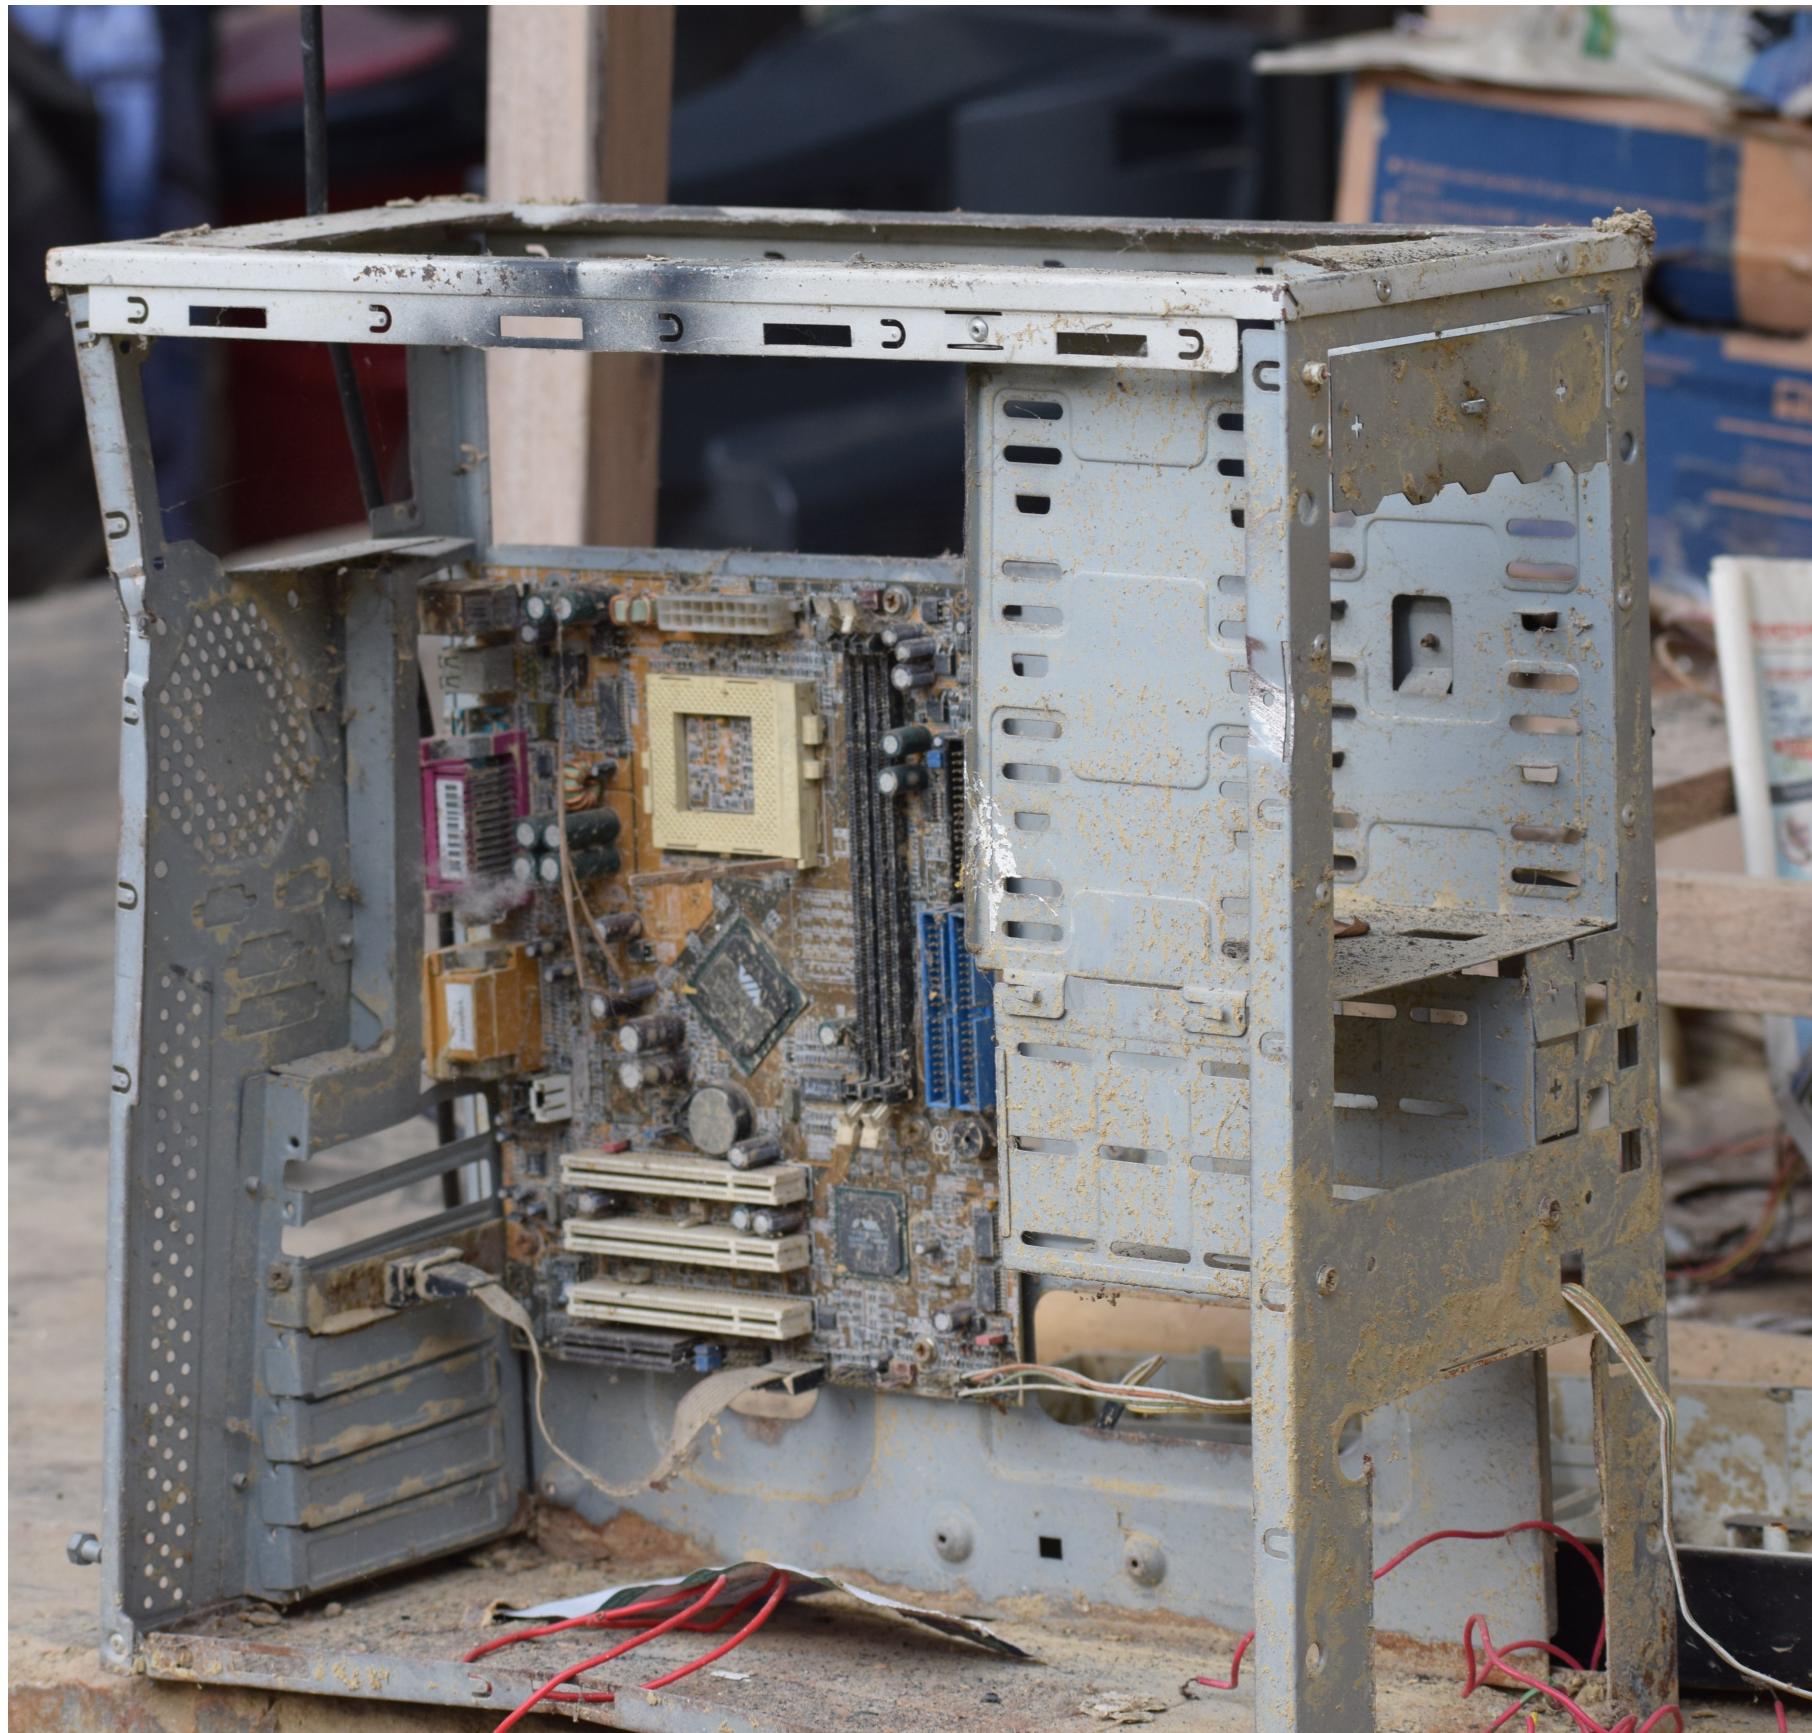
# What if *the network fails*?



# Shared Fate

Are you still there?

- Two methods/threads/processes running on the same computer generally have **shared fate** [Crashed/not]
- When two machines in a distributed system can't talk to each other, how do we know if the other is crashed?
- We call this a **split brain** problem



# CAP Theorem: Consistency or Availability

- Pick two of three:
  - Consistency: All nodes see the same data at the same time (strong consistency)
  - Availability: Individual node failures do not prevent survivors from continuing to operate
  - Partition tolerance: The system continues to operate despite message loss (from network and/or node failure)
    - Can't drop this for a DS - networks can always fail

# DNS uses a relaxed consistency model

- Recall - DNS requirements for updating/creating records:
  - Latency - OK if it takes minutes/hours for an update to take full effect
  - Performance - Expect far fewer writes than reads
- DNS does not attempt to provide strong consistency
- Updates are made in a best-effort approach
- OK for clients to see the “old” value for a short time after it was updated

# Distributed Software Engineering Abstractions

## Key Question: Consistency vs Availability

- Distributed system will never match exact semantics of non-distributed system
- For replication do we value more: guaranteed consistency (looks like a single machine) or guaranteed availability (sometimes read stale data)?
  - For a lock server?
  - For the order of tweets on twitter?
- For partitioning: Where can we draw the line? How do we determine how to organize the partitions?
- Next module: common architectures for distributed systems

# **Review: Learning Objectives for this Lesson**

**By the end of this lesson, you should be able to...**

- Describe 5 key goals of distributed systems
- Understand the fundamental constraints of distributed systems
- Understand the roles of replication and partitioning in the context of the DNS infrastructure